

CS231A

Computer Vision: From 3D Reconstruction to Recognition

Class Time

M-W; 11:30—12:50PM



CS231A

Instructors



Silvio Savarese

- ssilvio@stanford.edu
- Office: Gates Building, room: 154
- Office hour: Friday 2-3pm or by appointment



Jeanette Bohg

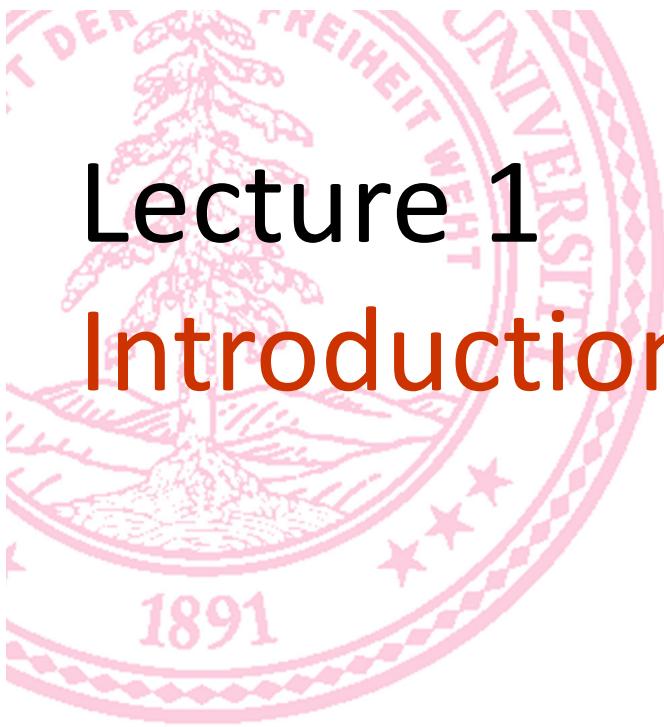
- bohg@stanford.edu
- Office: Gates Building, room: 140
- Office hour: Friday 1-2pm or by appointment

CAs:

- Andrey Kurenkov, Kuan Fang, Brent Yi, Krishnan Srinivasan

Lecture 1

Introduction



- An introduction to computer vision
- Course overview

AI is a propelling force of today's technology

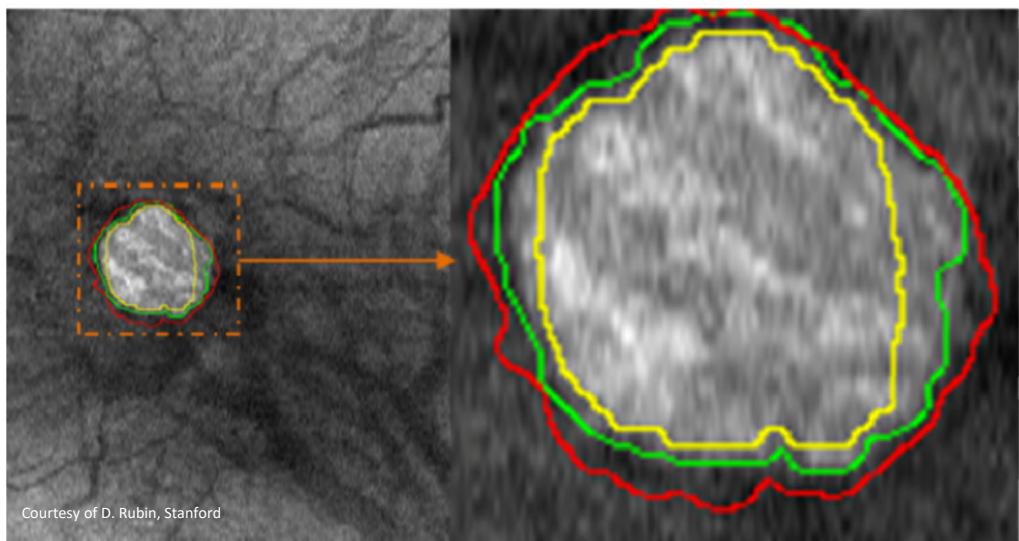


Smart Agriculture

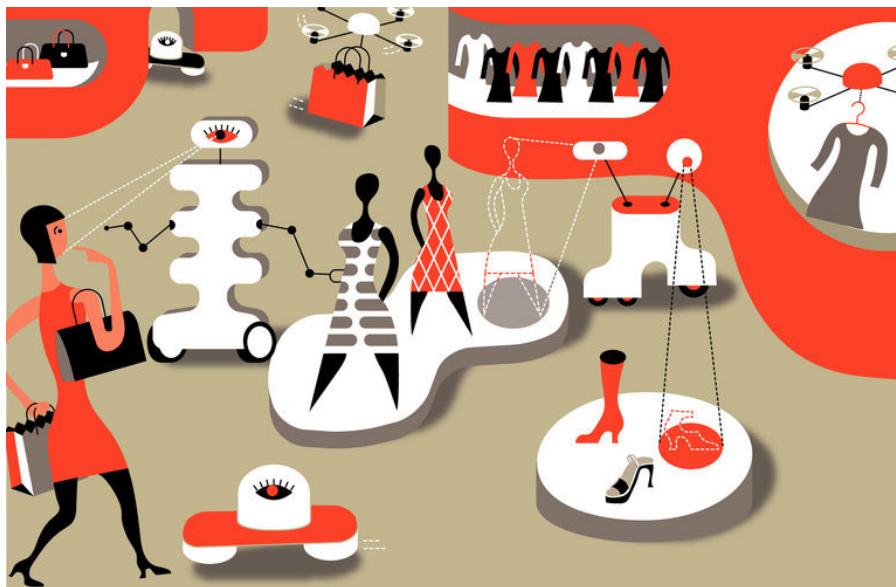
Courtesy of Agriculture Corner



Health care



Retail



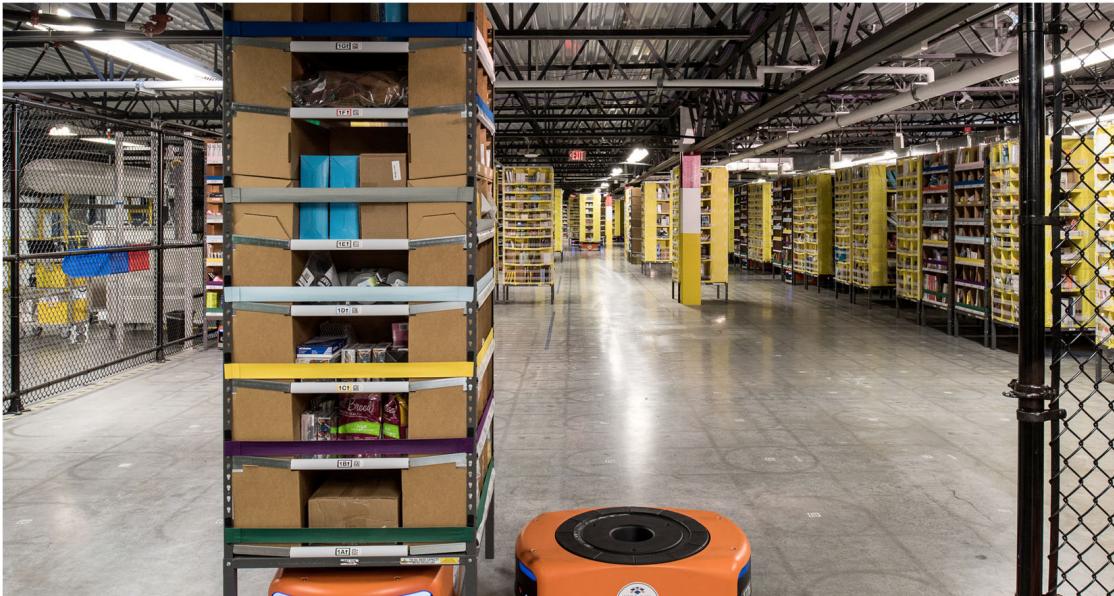
From Imagining the Retail Store of the Future - The New York Times, April 12, 2017

Manufacturing



Courtesy of ITRI, 2017

Transportation and Logistics



Construction Management



Why is this acceleration
happening now?

Enabling factors

- Big data



ImageNet, 2009
ShapeNet, 2015

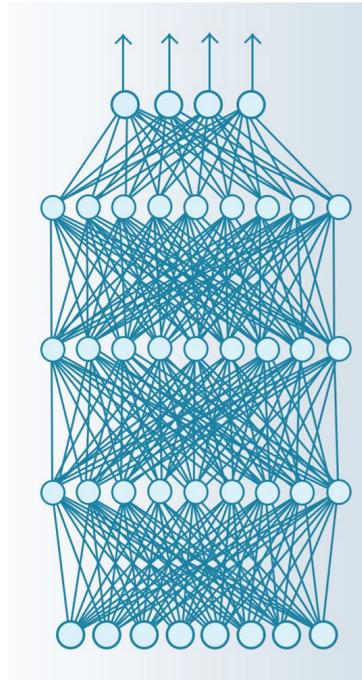
Enabling factors

- Big data
- Faster hardware

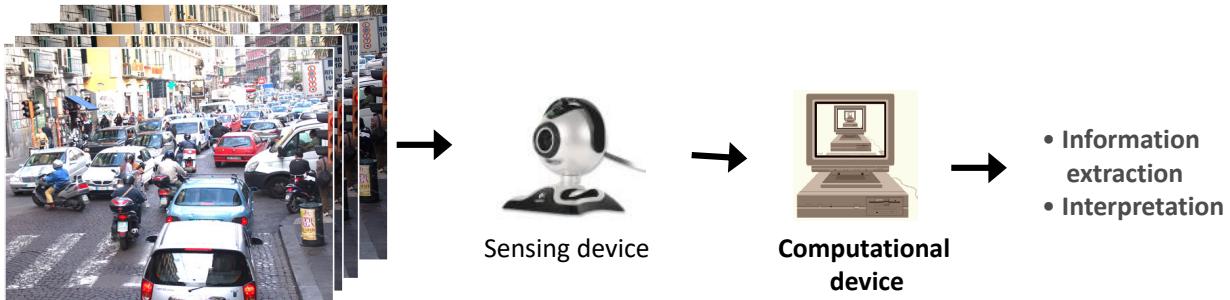


Enabling factors

- Big data
- Faster hardware
- New algorithms
 - Representation learning
 - Neural networks
 - Inject learning to deterministic reasoning



Computer vision



- 1. Information extraction:** features, 3D structure, motion flows, etc...
- 2. Interpretation:** recognize objects, scenes, actions, events

Major areas in Computer Vision



Space/Geometry

- Object shape recovery
- Depth estimation
- 3D scene reconstruction



Semantics/Learning

- Object detection and pose estimation
- Object tracking
- Scene understanding

Major areas in Computer Vision



Space/Geometry

- Object shape recovery
- Depth estimation
- 3D scene reconstruction



Semantics/Learning

- Object detection and pose estimation
- Object tracking
- Scene understanding

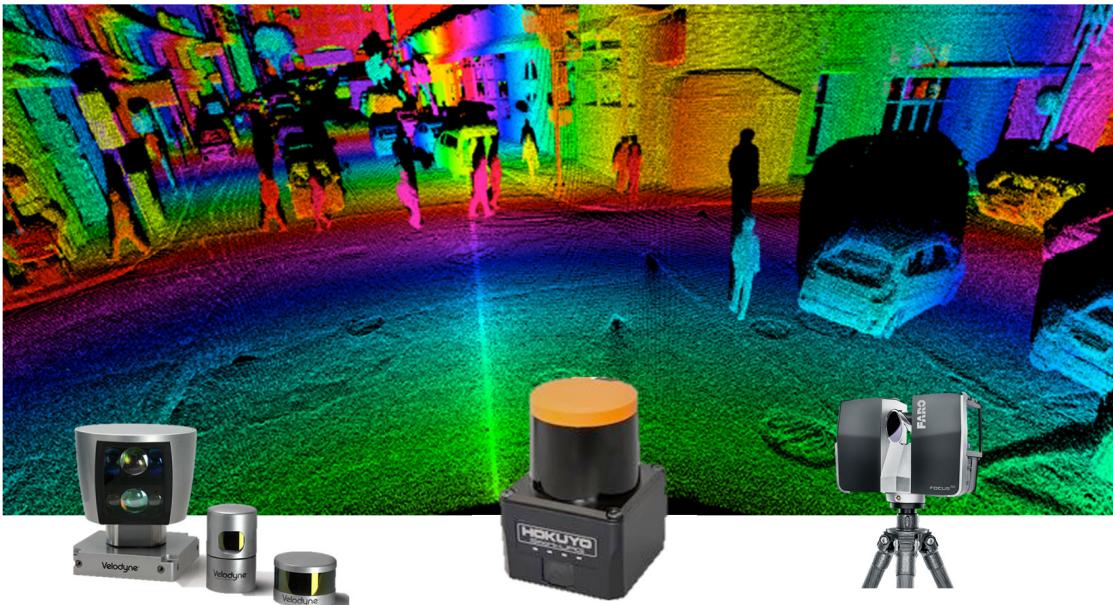
Recovering 3D models of the environments



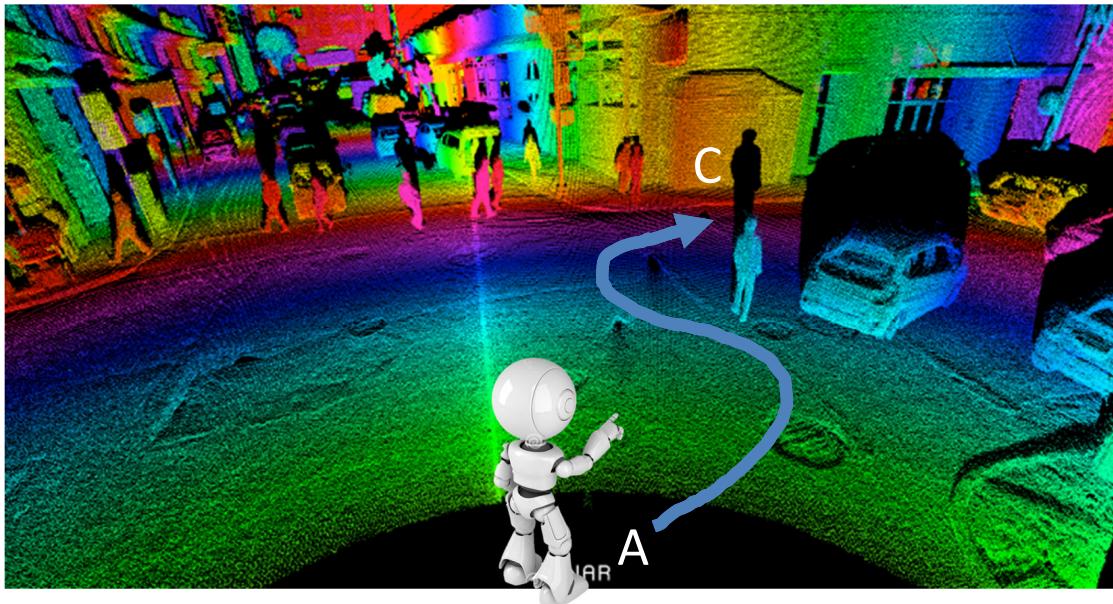
Armeni et al. 2016



Recovering 3D models of the environments



This is critical for autonomous
driving or navigation!



Major areas in Computer Vision



Space/Geometry

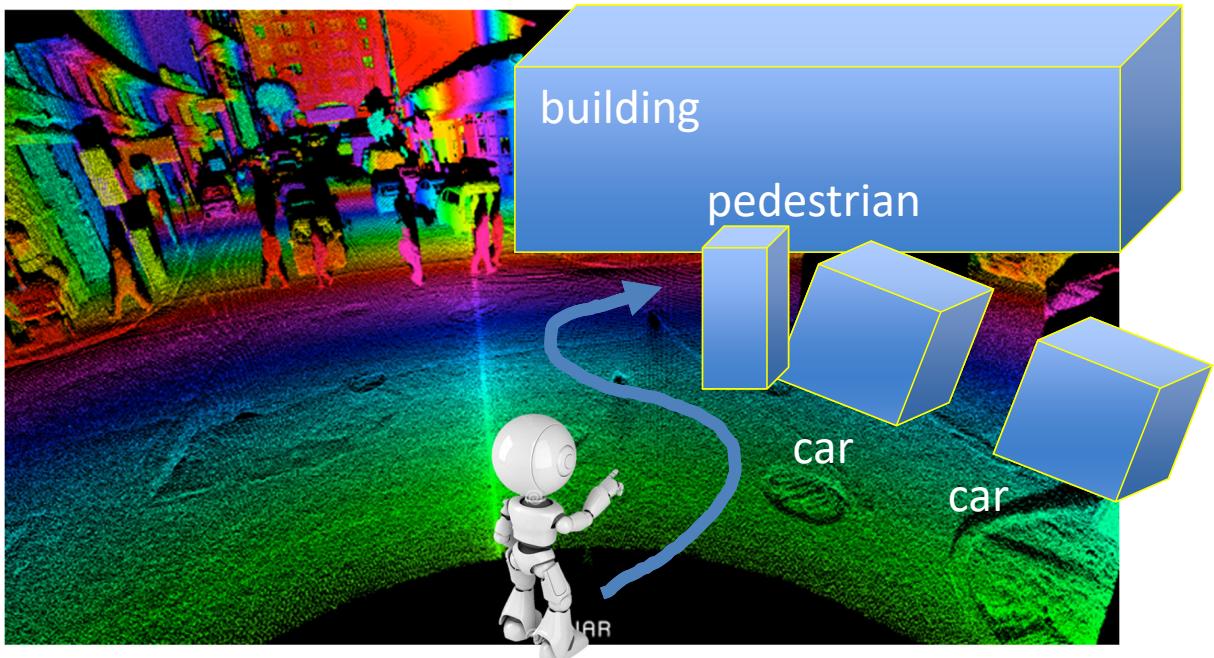
- Object shape recovery
- Depth estimation
- 3D scene reconstruction



Semantics/Learning

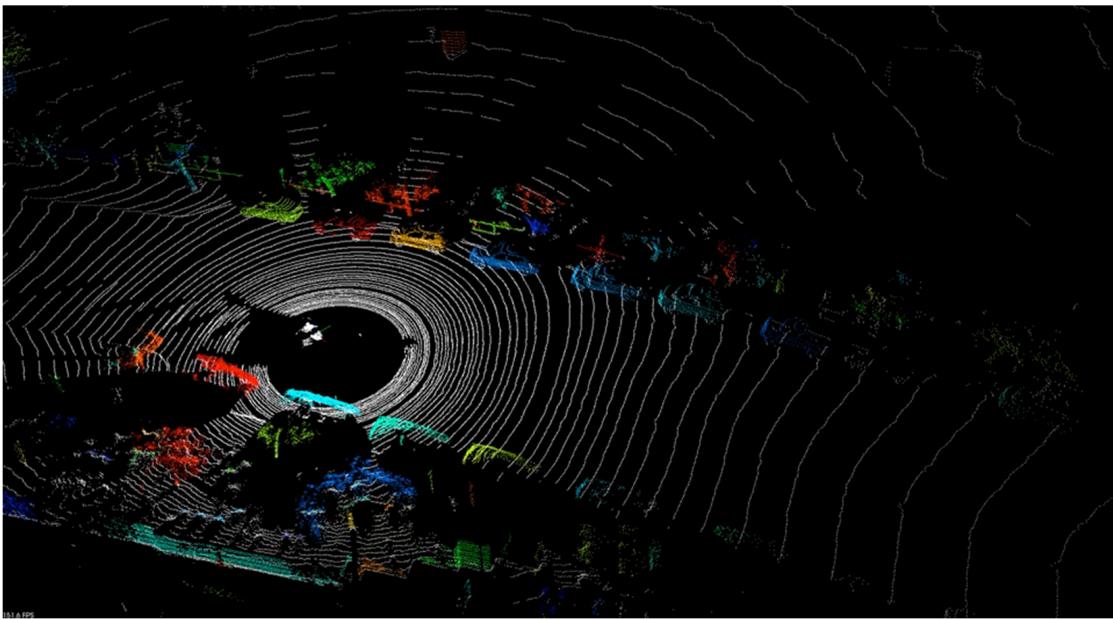
- Object detection and pose estimation
- Object tracking
- Scene understanding

Detecting and tracking objects in the environments



3D Scene Parsing

Held, Thrun, Savarese, 2016-206



Major areas in Computer Vision



Space/Geometry

- Object shape recovery
- Depth estimation
- 3D scene reconstruction



Semantics/Learning

- Object detection and pose estimation
- Object tracking
- Scene understanding

CS 231A course overview

1. Space/Geometry

Estimating spatial properties of objects and scene from images through geometrical methods

2. Semantics/Learning

CS 231A course overview

1. Space/Geometry

Estimating spatial properties of objects and scene from images through geometrical methods

2. Semantics/Learning

Estimating semantic and dynamic properties of scene elements from images through learning methods

CS 231A course overview

1. Space/Geometry

Estimating spatial properties of objects and scene from images through geometrical methods

2. Semantics/Learning

Estimating semantic and dynamic properties of scene elements from images through learning methods

Camera systems

Establish a mapping from 3D to 2D



How to calibrate a camera

Estimate camera parameters such pose or focal length



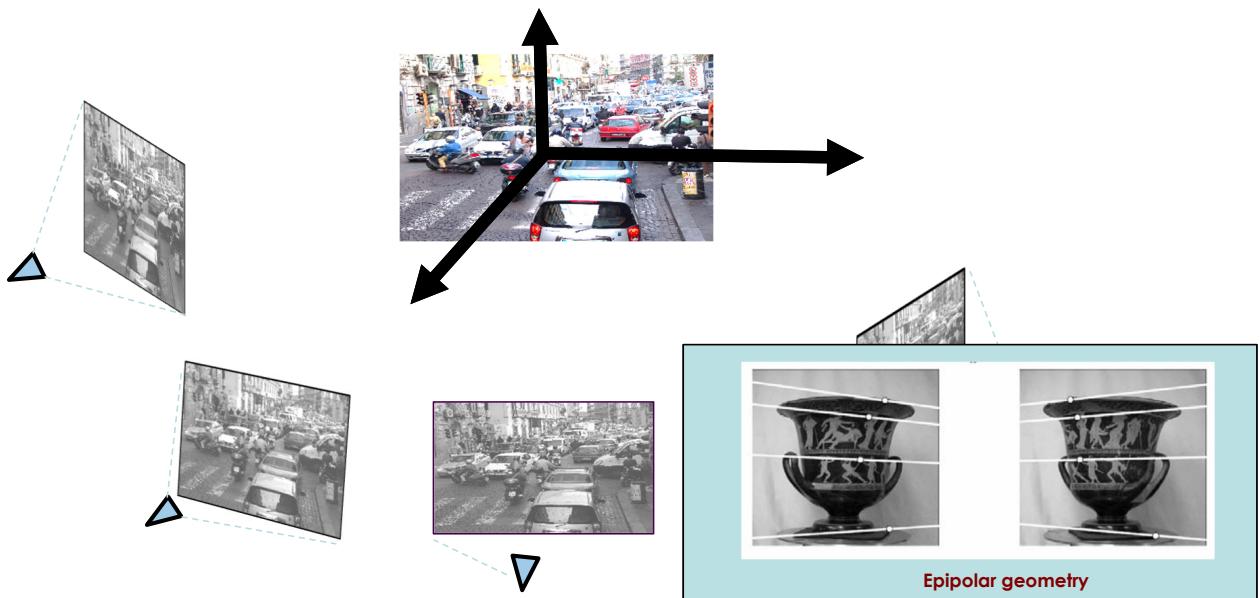
Single view metrology

Estimate 3D properties of the world from a single image

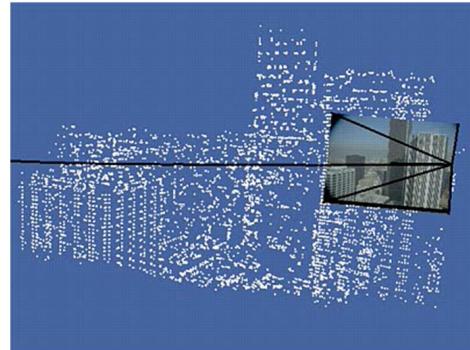


Multiple view geometry

Estimate 3D properties of the world from multiple views

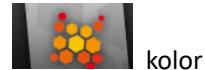
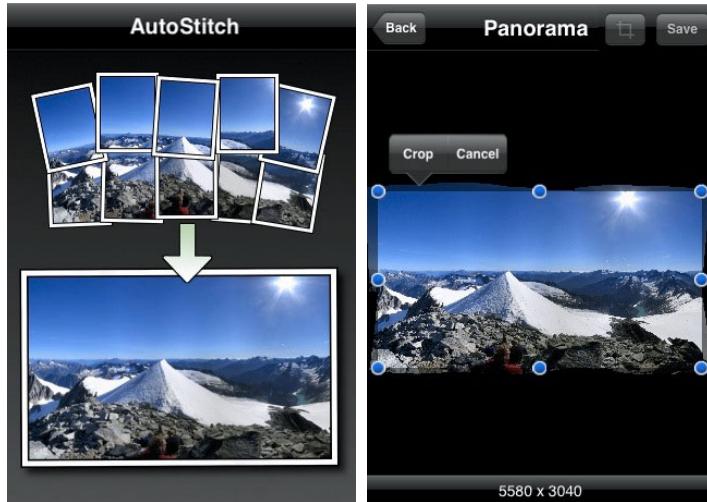


Structure from motion

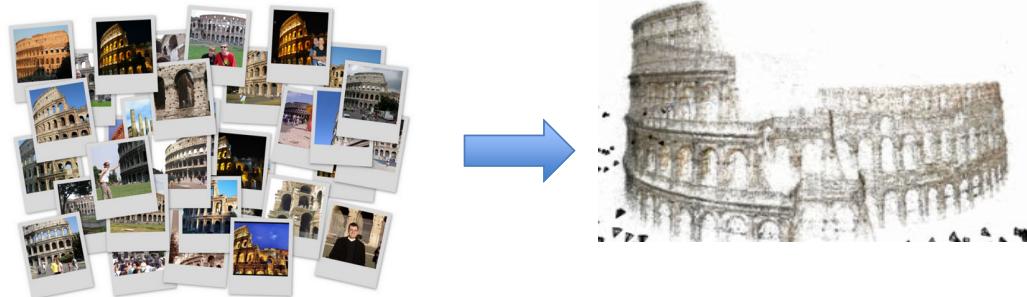


Courtesy of Oxford **Visual Geometry Group**

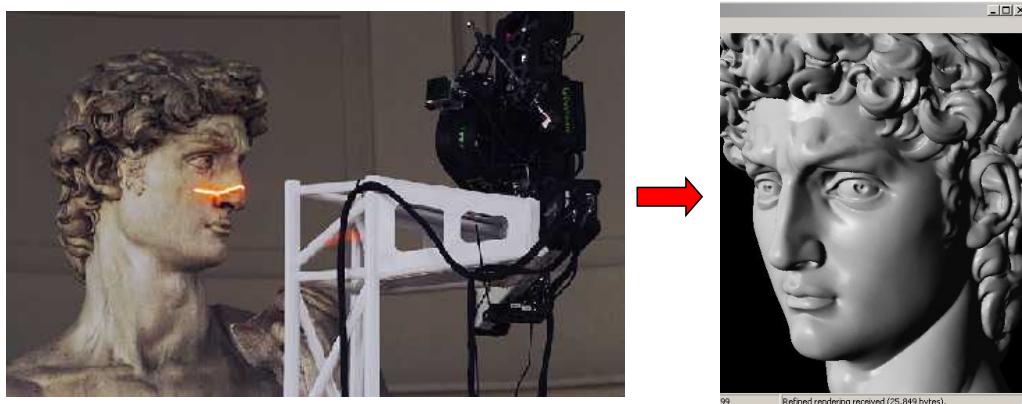
Panoramic Photography



3D Modeling of landmarks



Accurate 3D Object Prototyping



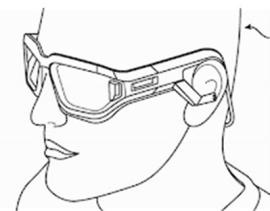
Scanning Michelangelo's "*The David*"

- [The Digital Michelangelo Project](#)
 - <http://graphics.stanford.edu/projects/mich/>
- 2 BILLION polygons, accuracy to .29mm

Augmented Reality



Mirriad
Advertising for the Skip Generation



- Magic leap
- Daqri
- Meta
- Etc...

CS 231A course overview

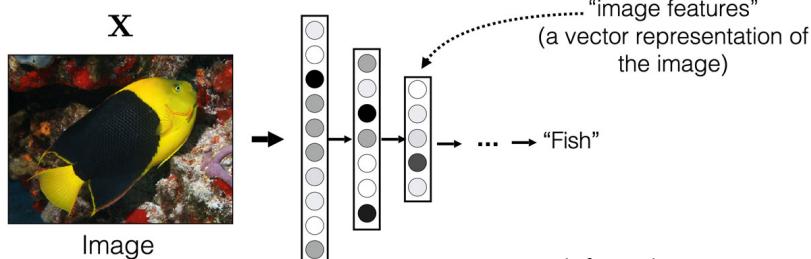
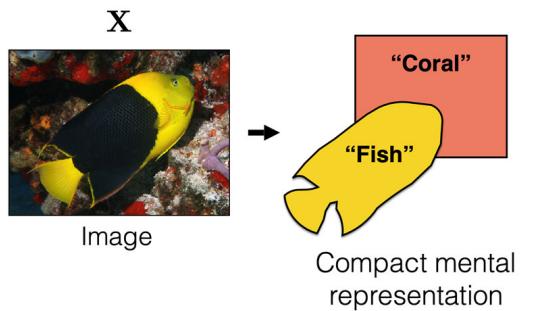
1. Space/Geometry

Estimating spatial properties of objects and scene from images through geometrical methods

2. Semantics/Learning

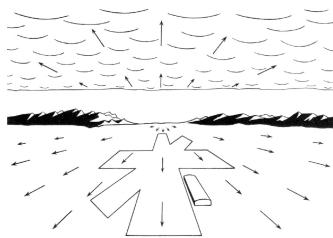
Estimating semantic and dynamic properties of scene elements from images through learning methods

Representations and Representation Learning



Example from Advances in Computer Vision – MIT – 6.869/6.819

Feature Tracking and Flow



J. J. Gibson, The Ecological Approach to Visual Perception



Lucas-Kanade Feature Tracking over multiple frames.
Picture adopted from OpenCV Webpage.

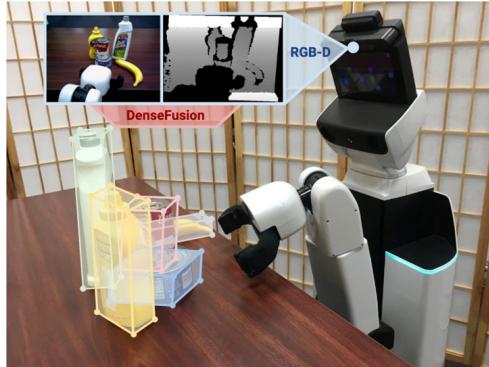


A Database and Evaluation Methodology for Optical Flow.
Baker et al. IJCV. 2011



A Primal-Dual Framework for Real-Time Dense RGB-D Scene Flow. Jaimez et al. ICRA, 2015.

Object Pose Estimation and Tracking

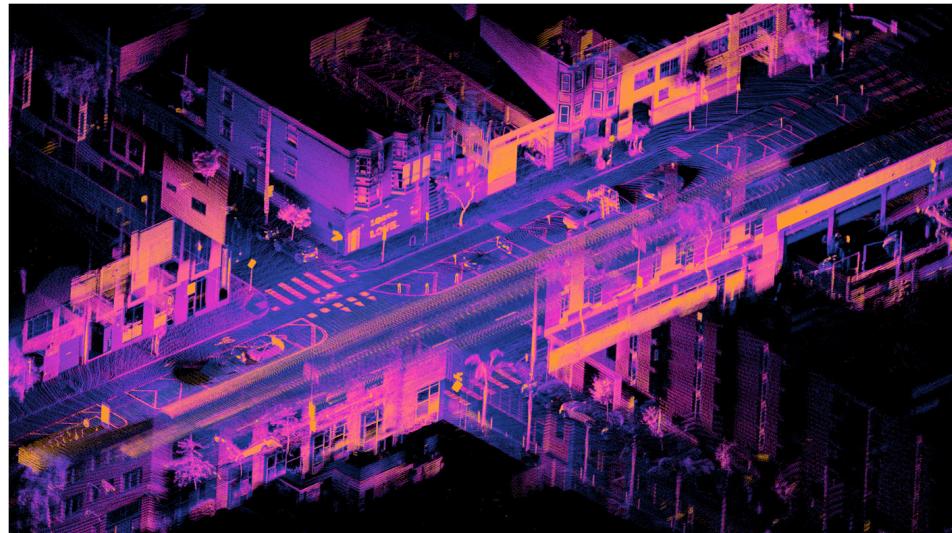


Z dqj hw#id##G hqvh#xvlrq#PG #R emfw#B rvh#ivwp dwrq# | #hudw#h#
G hqvh#xvlrq#P YSU 534<



P dpxhd# Ikvdk hw#ld##Surede#bw#R emfw#udfnbjj#kvbjj #B G hsvk#
F dp hud#EUR V#346

SLAM and Localization



Accumulated registered point cloud from lidar SLAM.

Autonomous navigation and safety

The image shows a screenshot of the Mobileye website. At the top, there are tabs for "manufacturer products" and "consumer products". Below this, a banner features the text "Our Vision. Your Safety." above an overhead view of a car with three camera systems highlighted: "rear looking camera", "forward looking camera", and "side looking camera".

Below the banner, there are three main sections:

- EyeQ Vision on a Chip**: Shows a close-up image of a chip labeled "EyeQ".
- Vision Applications**: Shows a person walking across a crosswalk with a yellow box highlighting the pedestrian.
- AWS Advance Warning System**: Shows a circular device with a display screen.

On the right side of the page, there is a "News" section with two items:

- Mobileye Advanced Technologies Power Volvo Cars World First Collision Warning With Auto Brake System
- Volvo: New Collision Warning with Auto Brake Helps Prevent Rear-end

Below the news is a "Events" section with two items:

- Mobileye at Equip Auto, Paris, France
- Mobileye at SEMA, Las Vegas, NV

[Mobileye](#): Vision systems in high-end BMW, GM, Volvo models
But also, Toyota, Google, Apple, Tesla, Nissan, Ford, etc....

Source: A. Shashua, S. Seitz

Personal robotics



jibo

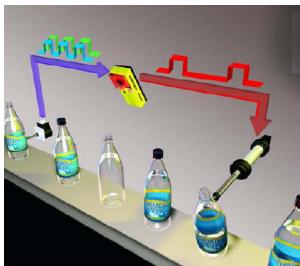
a

ECOVACS

FIVE
ELEMENTS
ROBOTICS

BLUE
FROG
robotics

More Applications



Factory inspection



Assistive technologies



Surveillance



Exploration and remote operations

Syllabus

January



Lecture	Topic
---------	-------

- 1 Introduction
- 2 Camera models
- 3 Camera calibration
- 4 Single view metrology
- 5 Epipolar geometry
- 6 Multi-view geometry

7 Structure from motion/ SLAM

8 Volumetric stereo

9 Fitting and Matching

10 Low Level Representations

11 Depth Estimation, Low Level Tracking

12 Optical and Scene Flow

13 6D pose Estimation, Object Tracking

14 Object Tracking Continued

15 SLAM

16 Guest

Geometry

← Proposal due

← Mid term

Learning

February



March



Project presentations

← Final projects

Prerequisites

- This course requires knowledge of linear algebra, probability, statistics, machine learning and computer vision, as well as decent programming skills.
- Though not an absolute requirement, it is encouraged and preferred that you have at least taken either CS221 or CS229 or CS131A or have equivalent knowledge.
- We will leverage concepts from low-level image processing (CS131A) (e.g., linear filters, edge detectors, corner detectors, etc...) and machine learning (CS229) (e.g., **SVM**, **basic Bayesian inference**, **clustering**, neural networks, etc...) which we won't cover in this class.
- We will provide links to background material related to CS131A and CS229 (or discuss during TA sessions) so students can refresh or study those topics if needed.

Text books

Required:

- [FP] D. A. Forsyth and J. Ponce. *Computer Vision: A Modern Approach* (2nd Edition). Prentice Hall, 2011.
- [HZ] R. Hartley and A. Zisserman. *Multiple View Geometry in Computer Vision*. Academic Press, 2002.

Recommended:

- R. Szeliski. *Computer Vision: Algorithms and Applications*. Springer, 2011.
- D. Hoiem and S. Savarese. *Representations and Techniques for 3D Object Recognition and Scene Interpretation*, Synthesis lecture on Artificial Intelligence and Machine Learning. Morgan Claypool Publishers, 2011
- Learning OpenCV, by Gary Bradski & Adrian Kaehler, O'Reilly Media, 2008.

Course assignments

- 1 warm up problem set (HW-0)
 - 4 problem sets
 - 1 mid-term exam
 - 1 project
-
- Look up class schedule for release and due dates.
 - Problems will be released through the [schedule page](#) and must be submitted through [Gradescope](#).

Midterm Exam

- The exam will be on 02/22. It will be released on Canvas and be available for 48 hours. You will have 2 hours to complete it once you start.
- You will be updated with more details, e.g., material to be covered, review sessions etc., as we approach the midterm.

Course Projects

- Replicate an interesting paper
 - Comparing different methods to a test bed
 - A new approach to an existing problem
 - Original research
-
- Write a 10-page paper summarizing your results
 - Release the final code
 - Give a final in-class presentation
 - SCPD students can send videos instead.
-
- We will introduce projects in 1-2 weeks
 - Important dates: look up class schedule

Course Projects

- Form your team:
 - 1-3 people
 - The larger is the team, the more work we expect from the team
 - Be nice to your partner: do you plan to drop the course?
- Evaluation
 - Quality of the project (including writing)
 - Final project in-class presentation (~ TBA minutes spotlight presentations)

Grading policy

- Homeworks: 37%
 - 1% for HW0
 - 9% for HW1, HW2, HW3, HW4 (each)
- Mid term exam: 20%
- Course project: 38%
 - Project proposal 1%
 - mid term progress report 5%
 - final report 25%
 - presentation 7%
- Attendance and class participation: 5%
 - Questions, answers, remarks, piazza posts,...
 - Class participation are waived for SCPD students. For the project presentation, SCPD students can send videos instead.

Grading policy (HWs)

- 25% will be deducted per day late.
- Two 48-hours one-time late submission “bonuses” are available; that is, you can use this bonus to submit your HW late after at most 48 hours. This is one time deal: After you use all your bonuses, you must adhere to the standard late submission policy.
- No exceptions will be made.

Grading policy (project)

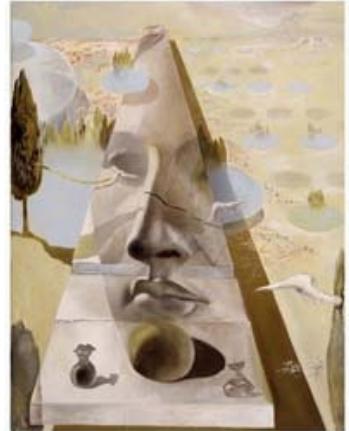
- If 1 day late, 25% off the grade for the project
- If 2 days late, 50% off the grade for the project
- Zero credits if more than 2 days
- No "late submission bonus" is allowed when submitting your progress report or project report



CS231

Introduction to

Computer Vision



Next lecture: Camera systems