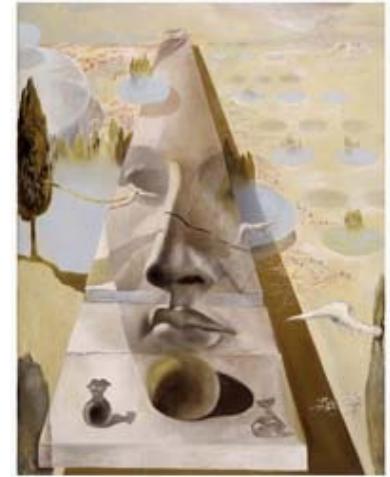


# CS231A

## Computer Vision: From 3D Reconstruction to Recognition



### Optical and Scene Flow

# What will you learn today?

Optical Flow

What is it and why do you care?

Assumptions

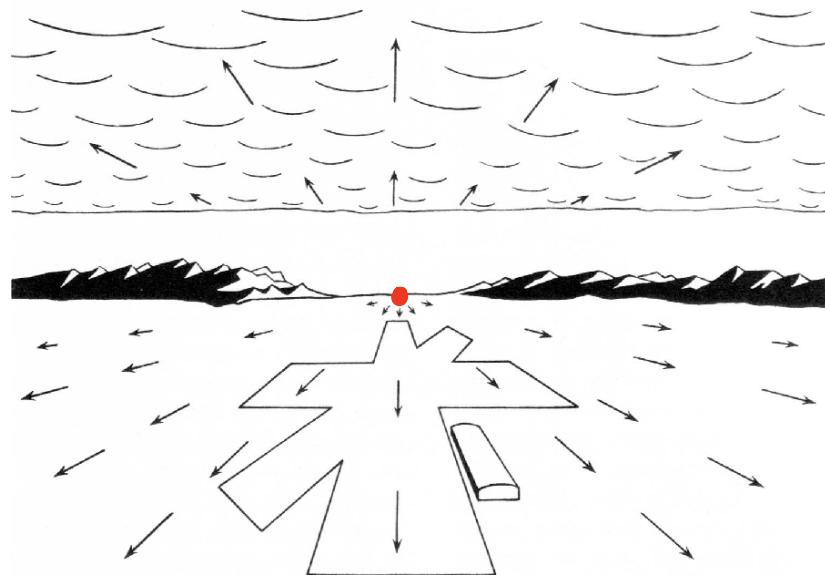
Formulating the optimization problem

Solving it

Scene Flow → 3D version of Optical flow

Learning-based Approaches to Estimating Motion

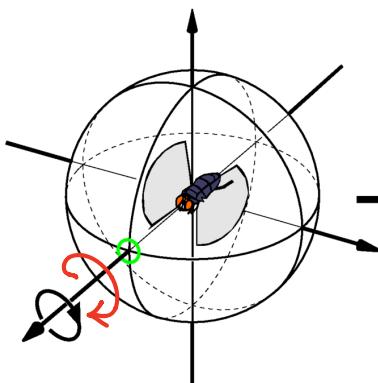
# Optical Flow - What is it?



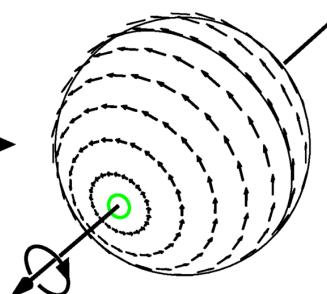
J. J. Gibson, *The Ecological Approach to Visual Perception*

# Optical Flow - What is it?

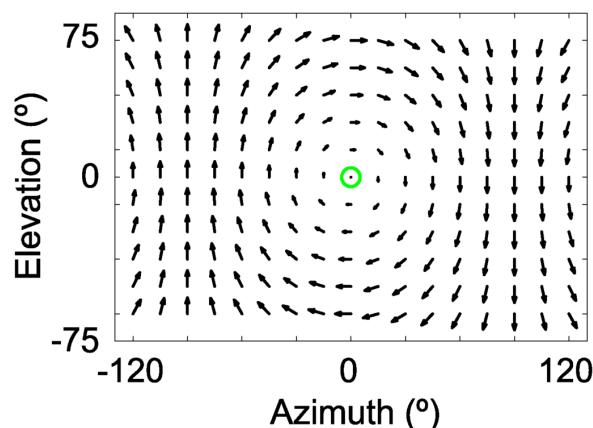
Rotation of observer



Optic flow  
(3D representation)

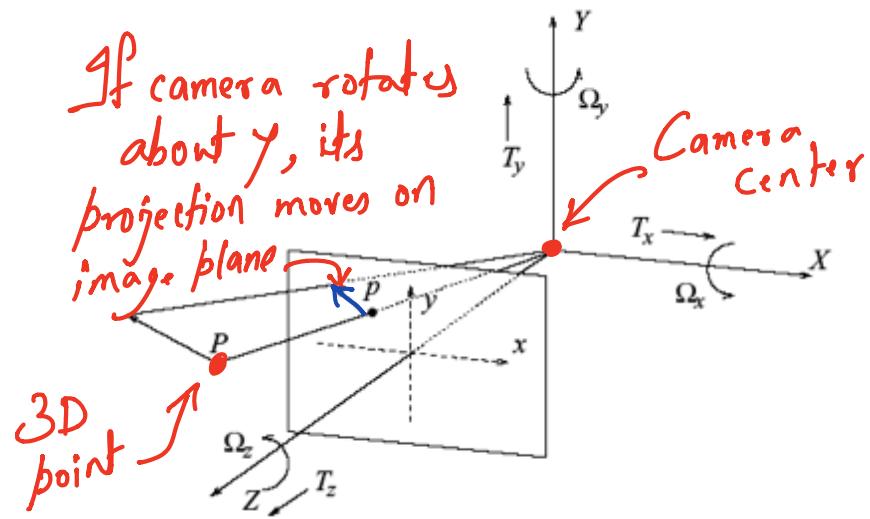


Optic flow  
(2D representation)



**Image Credit: Wikipedia. Optical Flow.**

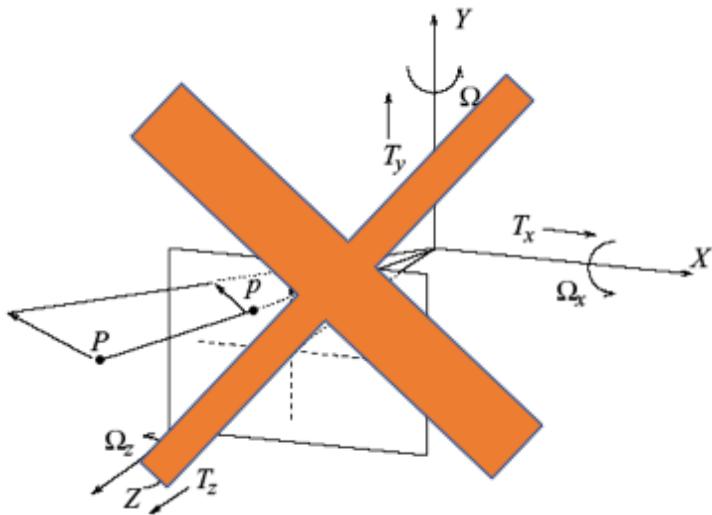
# Optical flow - What is it?



Motion field = 2D motion field representing the projection of the 3D motion of points in the scene onto the image plane.

B. Horn, Robot Vision, MIT Press

# Optical flow - What is it?



Optical flow = 2D velocity field describing the **apparent** motion in the images.

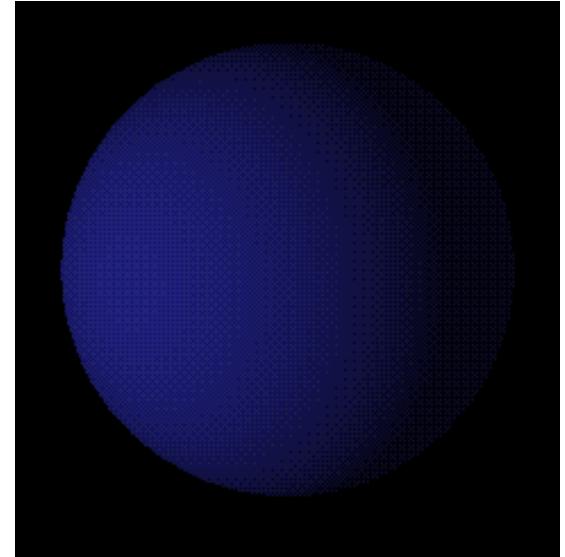
B. Horn, Robot Vision, MIT Press

# What is the motion field? What is the apparent motion?

Lambertian (matte) ball rotating in 3D

What does the 2D motion field look like?

What does the 2D optical flow field look like?



Slide Credit: Michael Black

Image source: <http://www.evl.uic.edu/aej/488/lecture12.html>

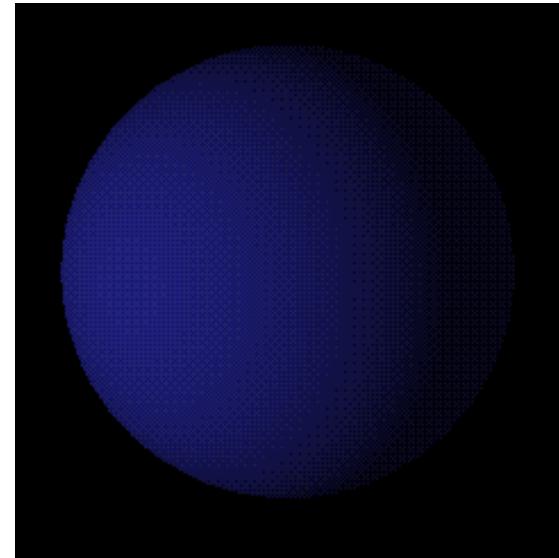
# What is the motion field? What is the apparent motion?

Stationary Lambertian (matte) ball

Moving Light Source

What does the 2D motion field look like?

What does the 2D optical flow field look like?



Slide Credit: Michael Black

Image source: <http://www.evl.uic.edu/aej/488/lecture12.html>

# Optical flow - What is it?

Motion Displacement of all image pixels

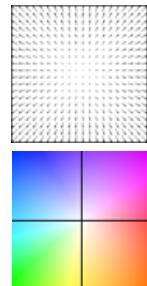


Image pixel value at time  $t$  and  
Location  $\mathbf{x} = (x, y)$ :  $I(x, y, t)$

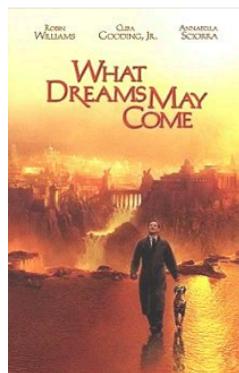
$u(x, y)$  horizontal component  
 $v(x, y)$  vertical component

Key

Slide Credit: Michael Black

# Optical Flow - What is it good for?

Painterly effect



Slide Credit: Michael Black

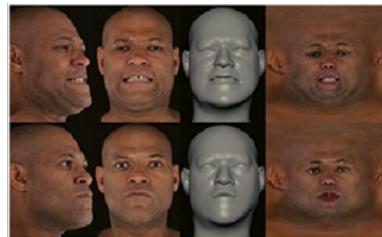
# Optical Flow - What is it good for?

Face morphing in matrix reloaded



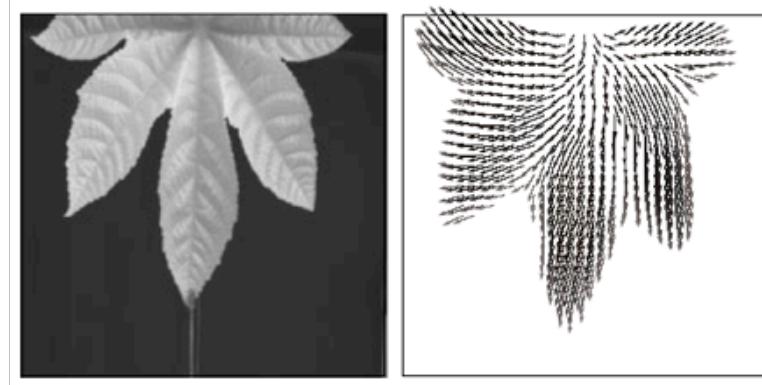
George Borshukov, Dan Piponi, Oystein Larsen, J.P.Lewis, Christina Tempelaar-Lietz  
ESC Entertainment

SIGGRAPH'03



Slide Credit: Michael Black

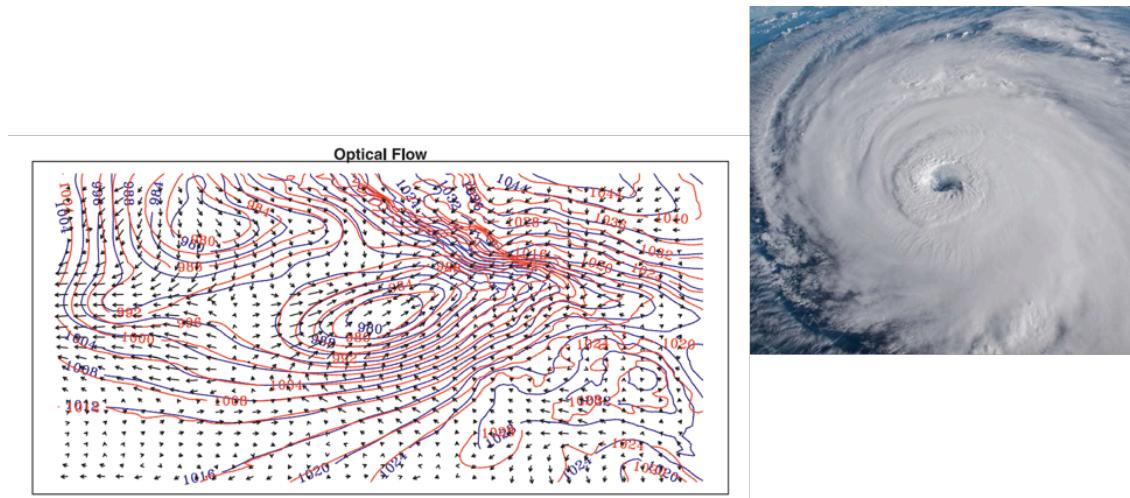
# Optical Flow - What is it good for?



Tensor-based Image Sequence Processing Techniques for the Study of  
Dynamical Processes, Haussecker, Spies, and Jahne, Proc. Intern. Symp.  
On Real-time Imaging and Dynamic Analysis: 704-711, 1998

Slide Credit: Michael Black

# Optical Flow - What is it good for?

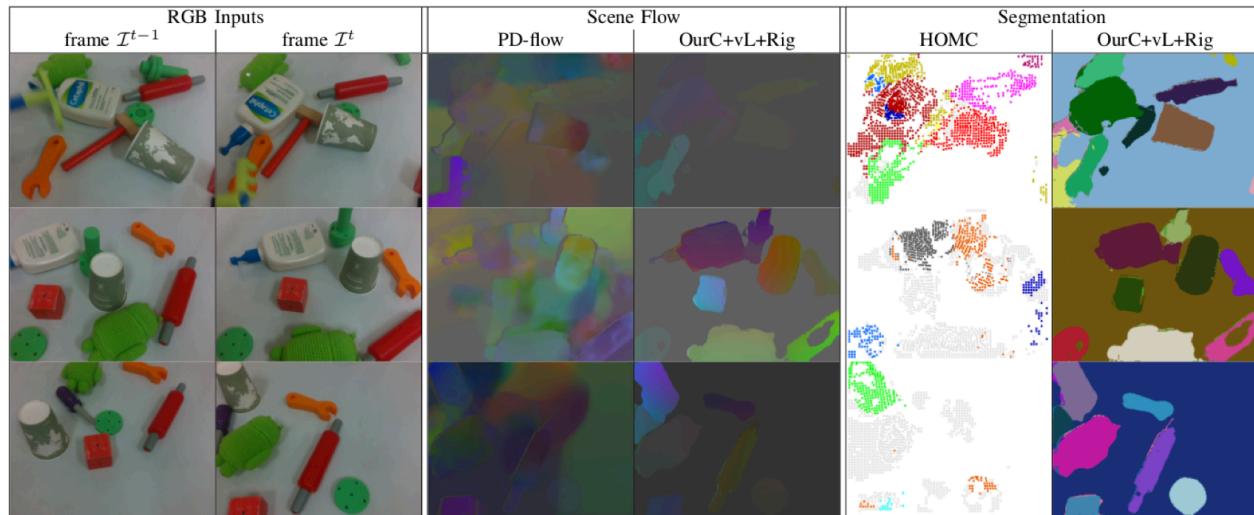


[Caren Marzban](#) and [Scott Sandgathe](#)

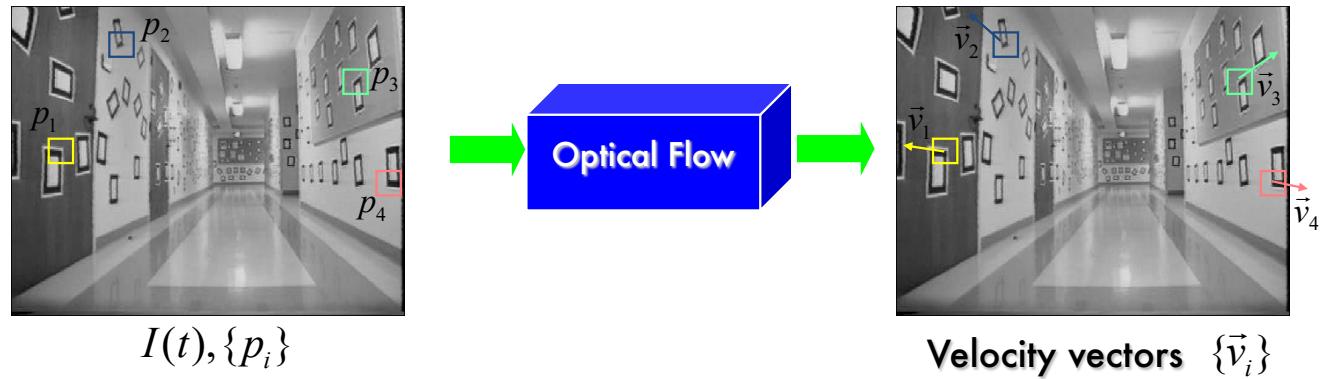
Optical Flow for Verification, Weather and Forecasting,  
Volume 25 No. 5, October 2010

Slide Credit: Michael Black

# Optical Flow - What is it good for?



# Optical Flow - What is it good for?



Slide Credit: CS223b – Sebastian Thrun

# Compute Optical Flow

## Goal

Compute the **apparent** 2D image motion of pixels from one image frame to the next in a video sequence.

# Compute (Sparse) Optical Flow



$I(t), \{p_i\}$



Velocity vectors  $\{\vec{v}_i\}$

# Simple KLT Tracker



Lucas  
An Iterative Image Registration Technique  
with an Application to Stereo Vision.



History of the  
Kanade-Lucas-Tomasi  
(KLT) Tracker

1981



Detection and Tracking of Feature Points.

1991

The original KLT algorithm



Good Features to Track.

1994

16-385 Computer Vision (Kris Kitani)

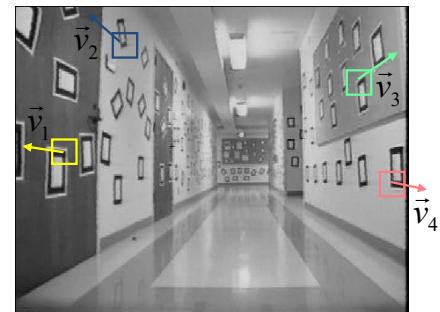
# Simple KLT Tracker

1. Find good points to track (Harris corners)
2. For each Harris corner compute motion (translation or affine) between consecutive frames
3. Link motion vector of successive frames to get a track for each Harris point
4. Introduce new Harris points by running detector every 10-15 frames
5. Track old and new corners using step 1-3

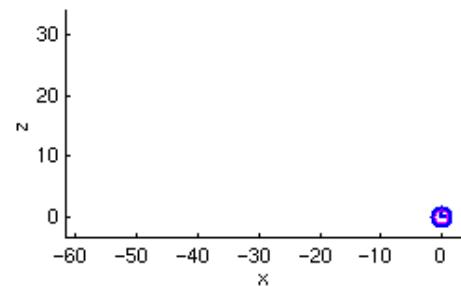
# Computing (Sparse) Optical Flow



$I(t), \{p_i\}$



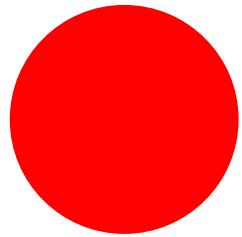
Velocity vectors  $\{\vec{v}_i\}$



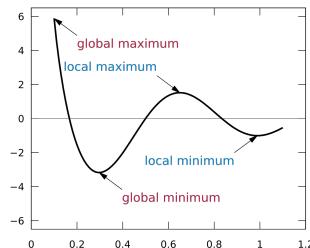
Jean-Yves Bouguet, [Ph.D.](#) CalTech

# Compute (Dense) Optical Flow

Step 1 - Assumptions



Step 2 - Objective Function

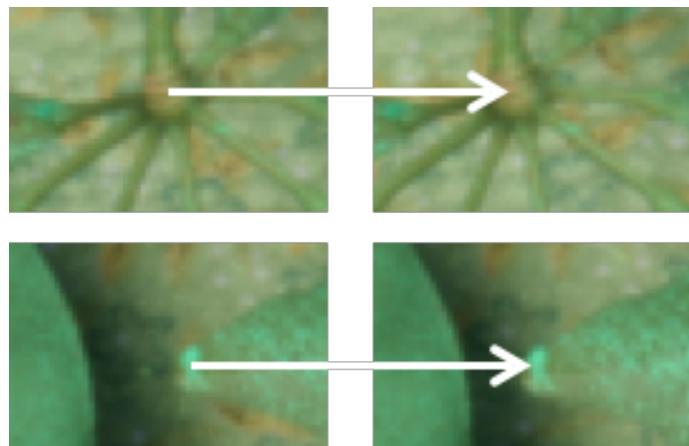


Source: Wikipedia.

Step 3 - Optimization



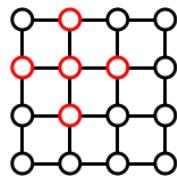
# Assumption 1 - Brightness Constancy



$$I(x + u, y + v, t + 1) = I(x, y, t)$$

Slide Credit: Michael Black

# Assumption 2 - Spatial Smoothness



Slide Credit: Michael Black

# Assumption 3 – Temporal Coherence

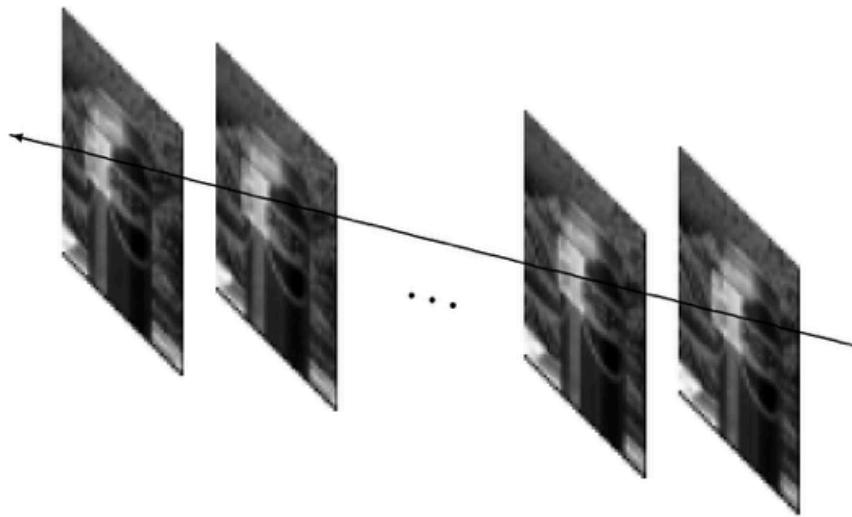
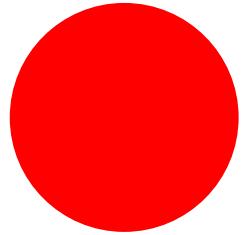


Figure 1.8: Temporal continuity assumption. A patch in the image is assumed to have the same motion (constant velocity, or acceleration) over time.

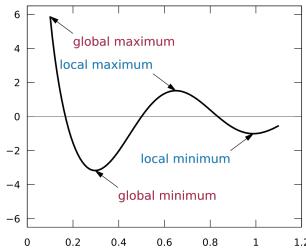
Slide Credit: Michael Black

# Compute Optical Flow

Step 1 - Assumptions



Step 2 - Objective Function



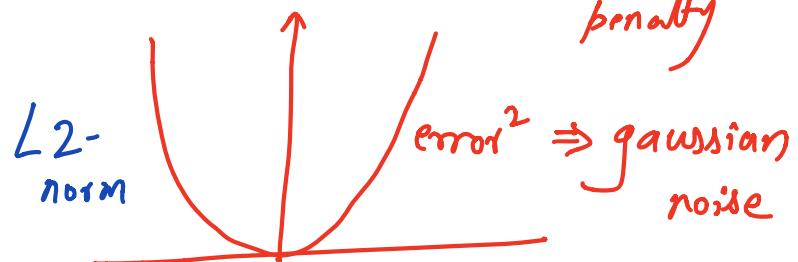
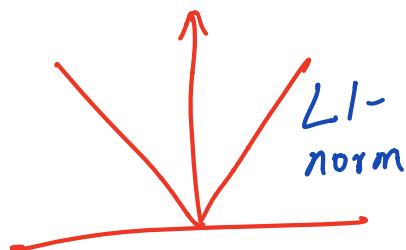
Source: Wikipedia.

$u, v \rightarrow$  optical field flow

# Objective Function

$$E_D(\mathbf{u}, \mathbf{v}) = \sum_{s \rightarrow \text{all pixels}} (I(x_s + u_s, y_s + v_s, t + 1) - I(x, y, t))^2$$

*data term*



New Assumption: Gaussian noise

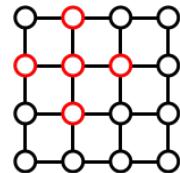
Huber loss  
(robust)

# Objective Function

*Spatial assumption*

$$E_S(\mathbf{u}, \mathbf{v}) = \sum_{n \in G(s)} (u_s - u_n)^2 + \sum_{n \in G(s)} (v_s - v_n)^2$$

*G(s) Neighborhood*



New Assumptions:

Flow field smooth

Gaussian Deviations

First order smoothness good enough

Flow derivative approximated by first differences

# Objective Function

$$E(u, v) = E_D(u, v) + \lambda E_S(u, v)$$

*relative weighting term*

$$E(u, v) = \sum_s (I(x_s + u_s, y_s + v_s, t + 1) - I(x, y, t))^2 + \lambda \left( \sum_{n \in G(s)} (u_s - u_n)^2 + \sum_{n \in G(s)} (v_s - v_n)^2 \right)$$

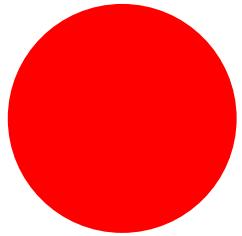
*Optimization variables*

*Need non-linear optimization*

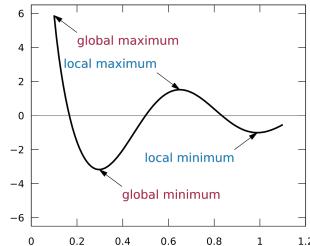


# Compute Optical Flow

Step 1 - Assumptions



Step 2 - Objective Function



Step 3 - Optimization



Source: Wikipedia.

# Linear Approximation

$$E(u, v) = \sum_s (I(x_s + u_s, y_s + v_s, t + 1) - I(x, y, t))^2 + \lambda (\sum_{n \in G(s)} (u_s - u_n)^2 + \sum_{n \in G(s)} (v_s - v_n)^2)$$

*Constraint equation for optical flow*

$u_s = dx, v_s = dy, dt = 1$

$$\cancel{I(x, y, t)} + \boxed{\cancel{I(x, y, t)} + dx \frac{\delta}{\delta x} I(x, y, t) + dy \frac{\delta}{\delta y} I(x, y, t) + dt \frac{\delta}{\delta t} I(x, y, t)} - I(x, y, t) = 0$$

*Partial derivative  
in  $x, y, f$  &  $t$  directions*

# Optical Flow Constraint Equation

Linearized loss function.

$$u \frac{\delta}{\delta x} I(x, y, t) + v \frac{\delta}{\delta y} I(x, y, t) + \frac{\delta}{\delta t} I(x, y, t) = 0$$

$$I_x u + I_y v + I_t = 0 \quad \rightarrow \text{Constraint equation at every pixel}$$

New Assumptions:

Flow is small

Image is differentiable

First order Taylor series is a good approximation

# Optical Flow Constraint Equation

## Notation

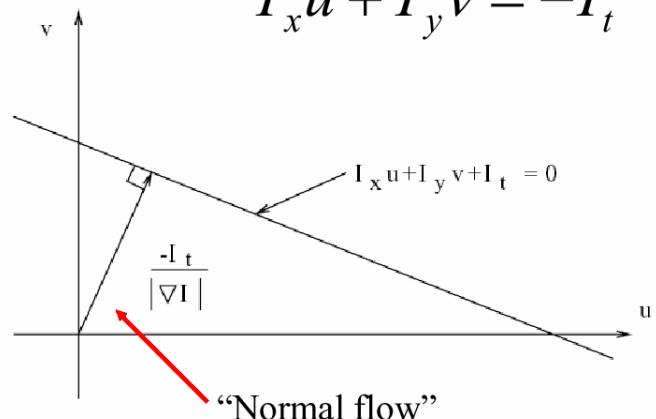
$$I_x u + I_y v + I_t = 0$$

$$\nabla I^T \mathbf{u} = -I_t$$

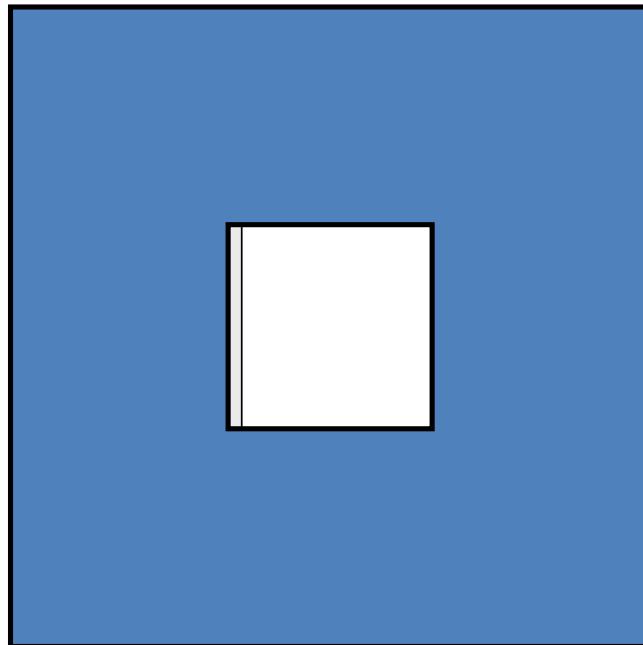
$$\mathbf{u} = \begin{bmatrix} u \\ v \end{bmatrix} \quad \nabla I = \begin{bmatrix} I_x \\ I_y \end{bmatrix}$$

At a single image pixel, we get a line:

$$I_x u + I_y v + I_t = 0$$

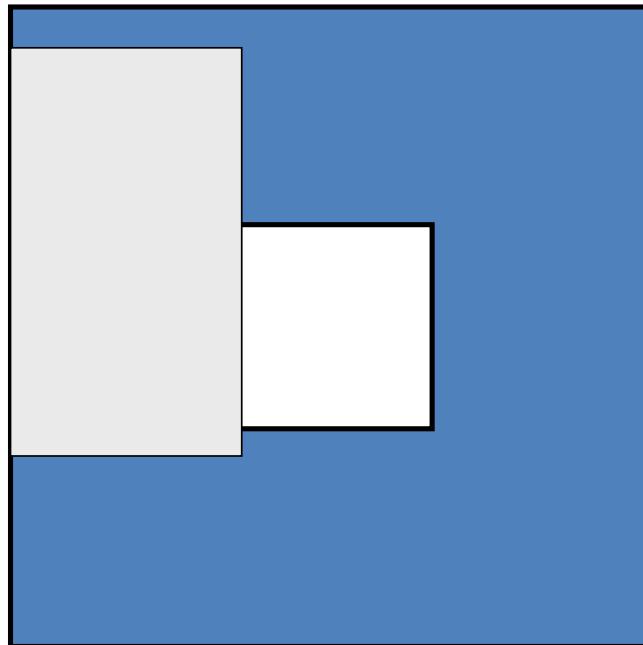


# Aperture Problem



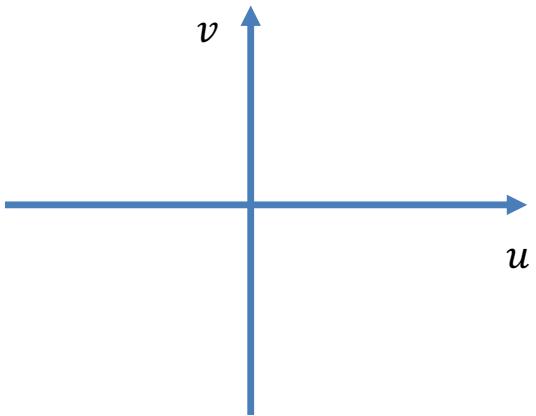
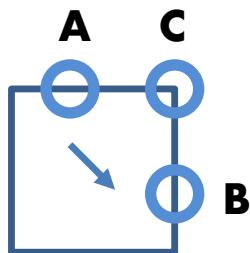
Slide Credit: CS223b – Sebastian Thrun

# Aperture Problem

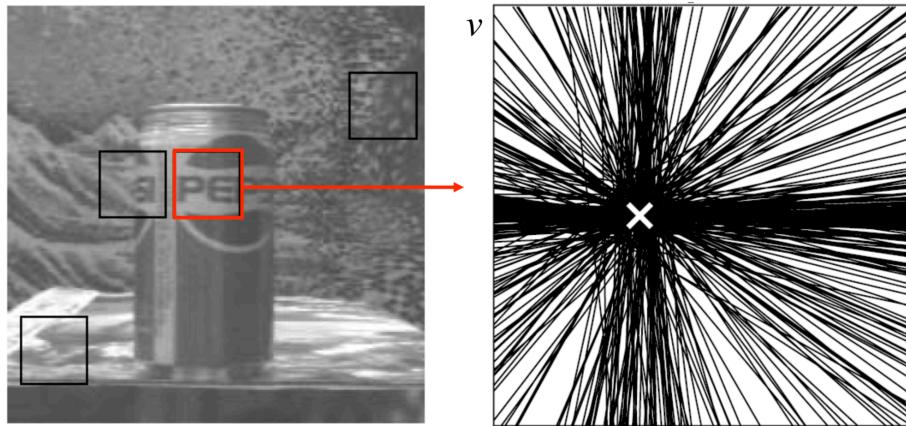


Slide Credit: CS223b – Sebastian Thrun

# What are the constraint lines?



# Multiple Constraints



Each pixel gives us a constraint:  $I_x u + I_y v = -I_t$

Slide Credit: Michael Black

# How do we solve this optimization problem?

$$E(u, v) = \sum_{x, y \in R} (I_x(x, y, t)u + I_y(x, y, t)v + I_t(x, y, t))^2$$

$\textcircled{R}$  → neighborhood

$$\frac{\partial E}{\partial u} = \sum_R (I_x u + I_y v + I_t) I_x = 0$$

$$\frac{\partial E}{\partial v} = \sum_R (I_x u + I_y v + I_t) I_y = 0$$

# How do we solve this optimization problem?

$$\left[ \sum_R I_x^2 \right] u + \left[ \sum_R I_x I_y \right] v = - \sum_R I_x I_t$$

$$\left[ \sum_R I_x I_y \right] u + \left[ \sum_R I_y^2 \right] v = - \sum_R I_y I_t$$

$$\begin{bmatrix} \sum R I_x^2 & \sum R I_x I_y \\ \sum R I_y I_x & \sum R I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} - \sum R I_x I_t \\ - \sum R I_y I_t \end{bmatrix}$$

# How do we solve this optimization problem?

$$\begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_y I_x & \sum I_y^2 \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix} = \begin{bmatrix} -\sum I_x I_t \\ -\sum I_y I_t \end{bmatrix}$$

$$\mathbf{A}\mathbf{u} = \mathbf{b}$$

# How do we solve this optimization problem?

If  $A$  is invertible

$$A^{-1}A\mathbf{u} = A^{-1}\mathbf{b}$$

$$\mathbf{u} = A^{-1}\mathbf{b}$$

$$A\mathbf{u} = \mathbf{b}$$

$$\mathbf{u} = \begin{bmatrix} u \\ v \end{bmatrix} \quad \nabla I = \begin{bmatrix} I_x \\ I_y \end{bmatrix}$$

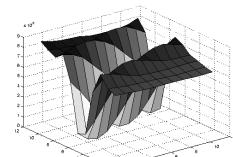
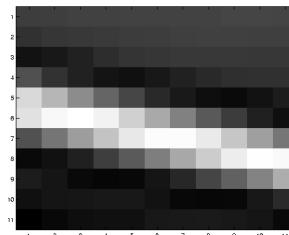
$$A^T A\mathbf{u} = A^T \mathbf{b}$$

$$\mathbf{u} = (A^T A)^{-1} A^T \mathbf{b}$$

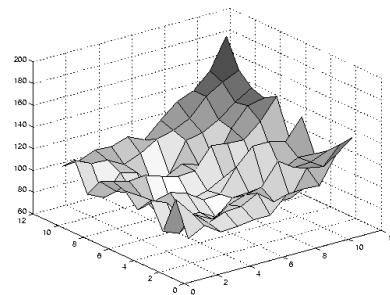
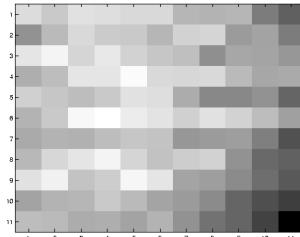
$$\mathbf{u} = - \left( \sum \nabla I \nabla I^T \right)^{-1} \sum \nabla I I_t$$

Use pseudo-inverse  
if not invertible

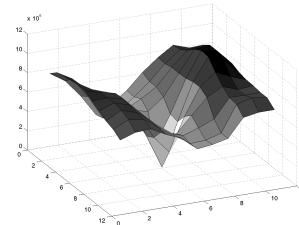
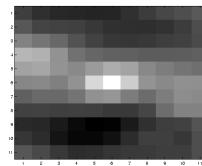
# Image Gradient Examples - Edge



# Image Gradient Examples – Low texture



# Image Gradient Examples – Low texture



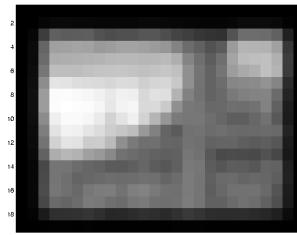
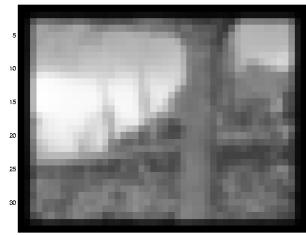
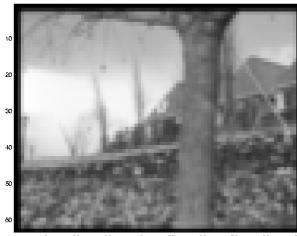
# Bag of tricks

Small motion assumption



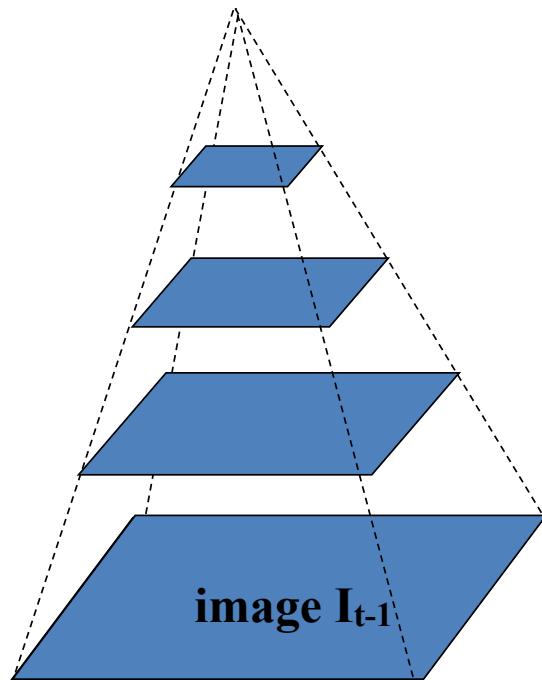
# Bag of tricks

## Reduce Resolution



\* From Khurram Hassan-Shafique CAP5415 Computer Vision 2003

# Spatial Pyramids



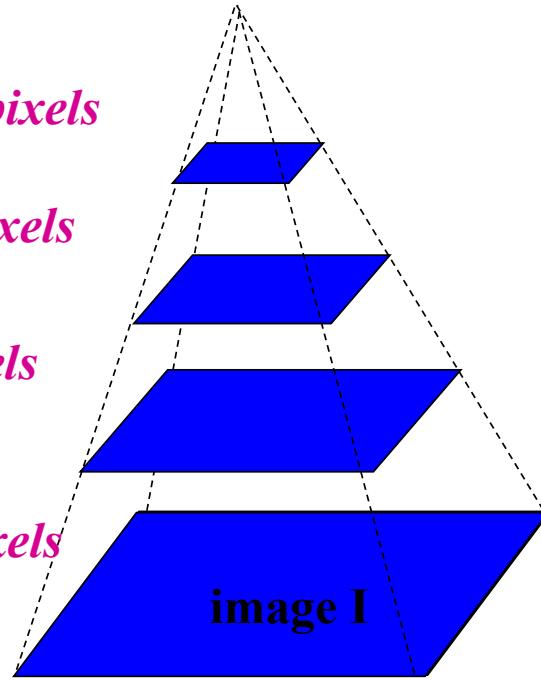
Gaussian pyramid of image  $I_{t-1}$

$u=1.25 \text{ pixels}$

$u=2.5 \text{ pixels}$

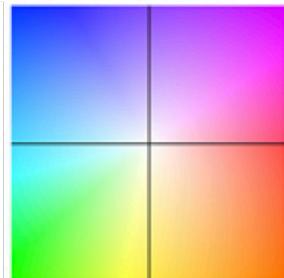
$u=5 \text{ pixels}$

$u=10 \text{ pixels}$

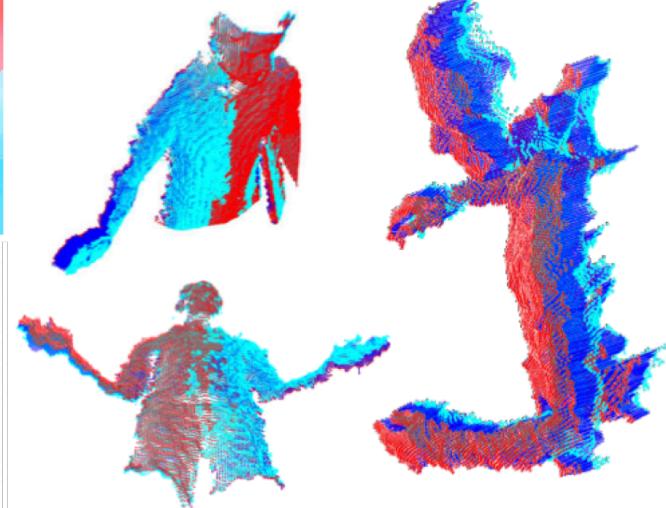


Gaussian pyramid of image  $I$

# Scene Flow



A Database and Evaluation Methodology for Optical Flow.  
Baker et al. IJCV, 2011



A Primal-Dual Framework for Real-Time Dense RGB-D Scene Flow. Jaimez et al. ICRA, 2015.

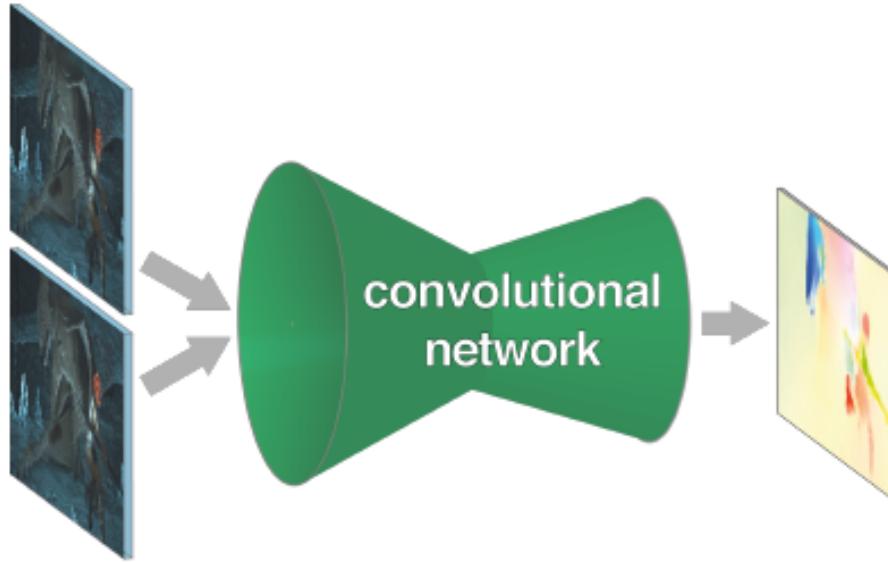
# What are the main challenges with this traditional formulation?

- Assumptions
  - Brightness constancy
  - Small motion
  - Etc
- Occlusions
- Large motion

# Learning-based approaches

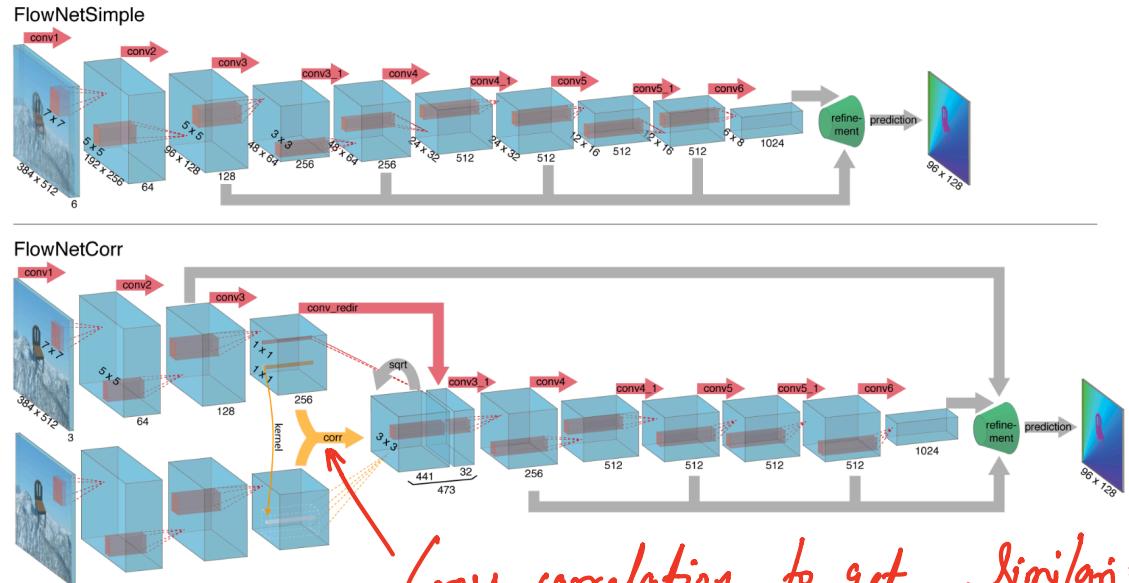
- Since 2015
- Availability of data

# FlowNet - Learning Optical Flow with Convolutional Networks



**Alexey Dosovitskiy, Philipp Fischer, Eddy Ilg, P. Häusser, C. Hazırbaş, V. Golkov, P. Smagt, D. Cremers, Thomas Brox. IEEE International Conference on Computer Vision (ICCV), 2015**

# FlowNet - Learning Optical Flow with Convolutional Networks

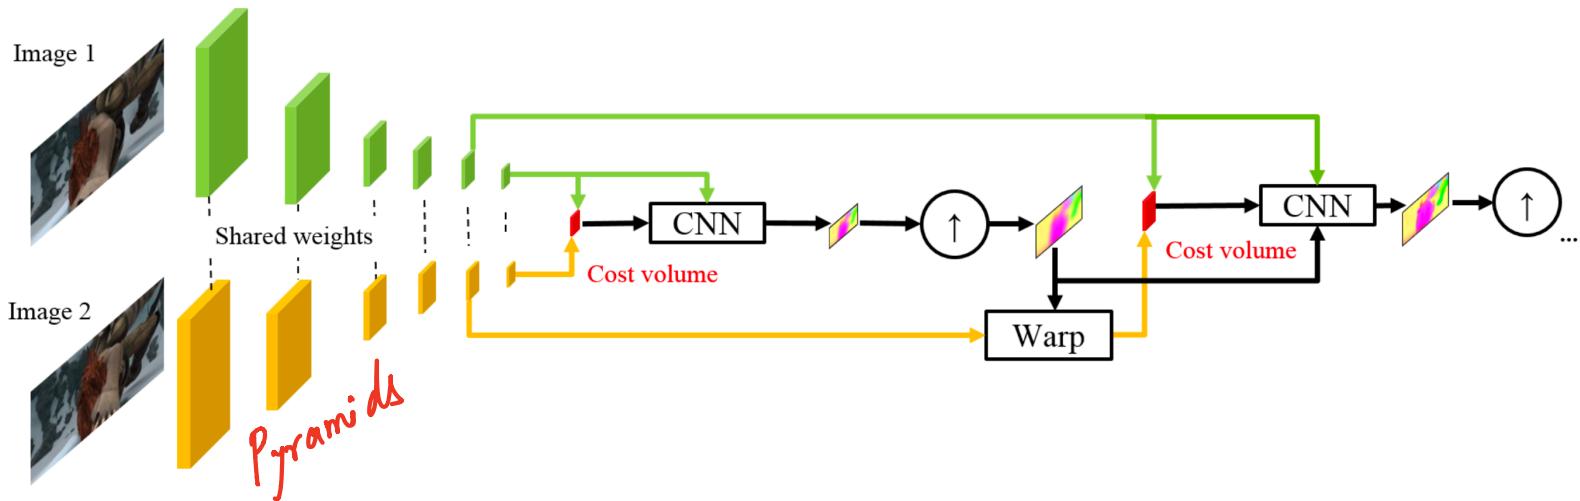


Supervised Learning

*Cross-correlation to get similarity between all patches in one image with all patches in the second image.*

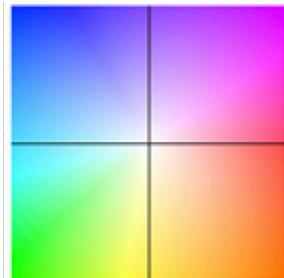
# Supervised Optical Flow using traditional principles

Pyramidal Processing, Image Warping & Cost volume processing

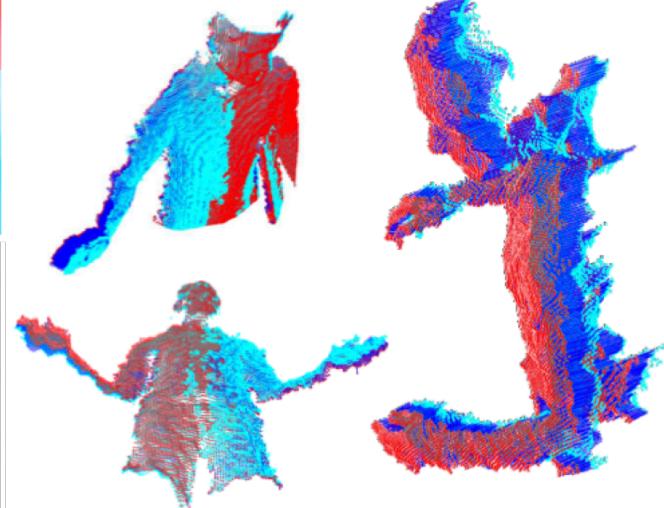


Deqing Sun, Xiaodong Yang, Ming-Yu Liu, and Jan Kautz. "PWC-Net: CNNs for Optical Flow Using Pyramid, Warping, and Cost Volume." CVPR 2018

# Scene Flow

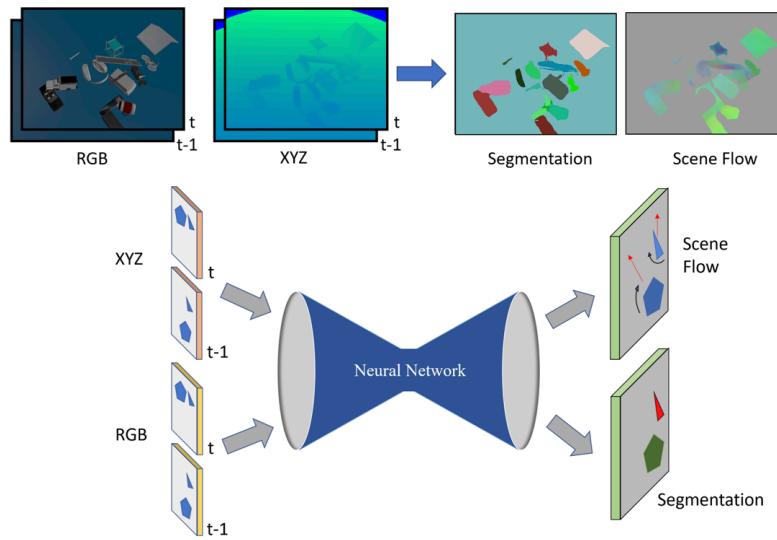


A Database and Evaluation Methodology for Optical Flow.  
Baker et al. IJCV. 2011



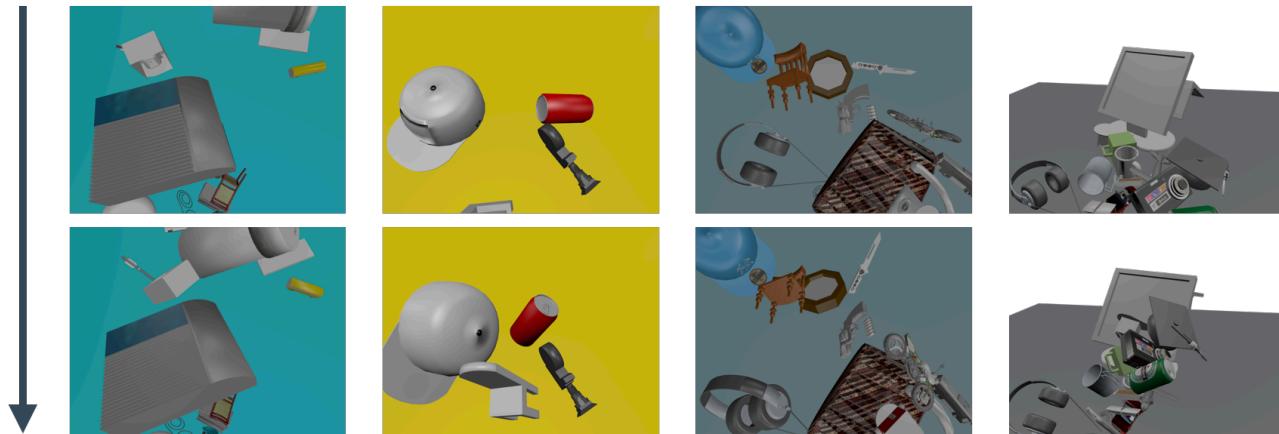
A Primal-Dual Framework for Real-Time Dense RGB-D Scene Flow. Jaimez et al. ICRA, 2015.

# Motion-based Object Segmentation based on Dense RGB-D Scene Flow



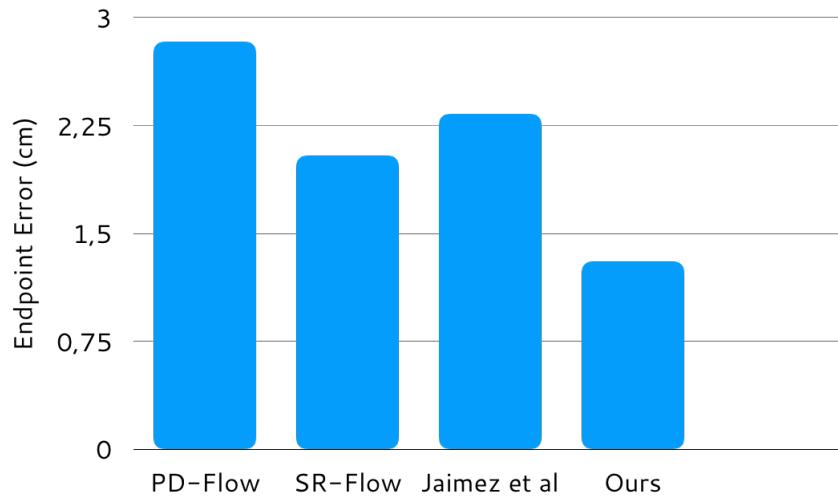
# Motion-based Object Segmentation based on Dense RGB-D Scene Flow

Time

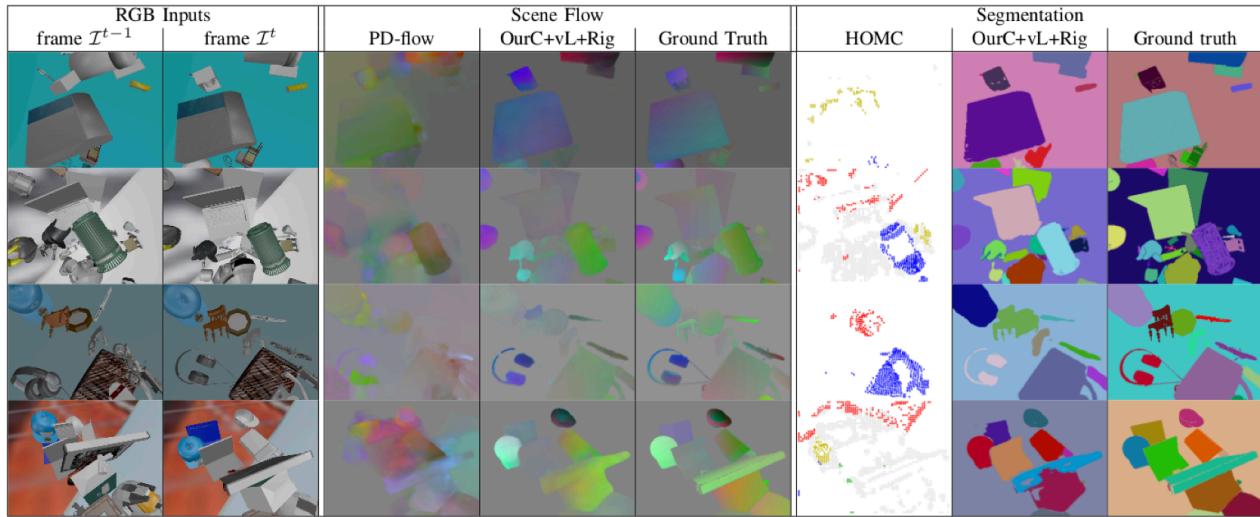


Motion-based Object Segmentation based on Dense RGB-D Scene Flow. Shao et al. RAL + IROS. 2018. Pre-print on arXiv.

# Motion-based Object Segmentation based on Dense RGB-D Scene Flow

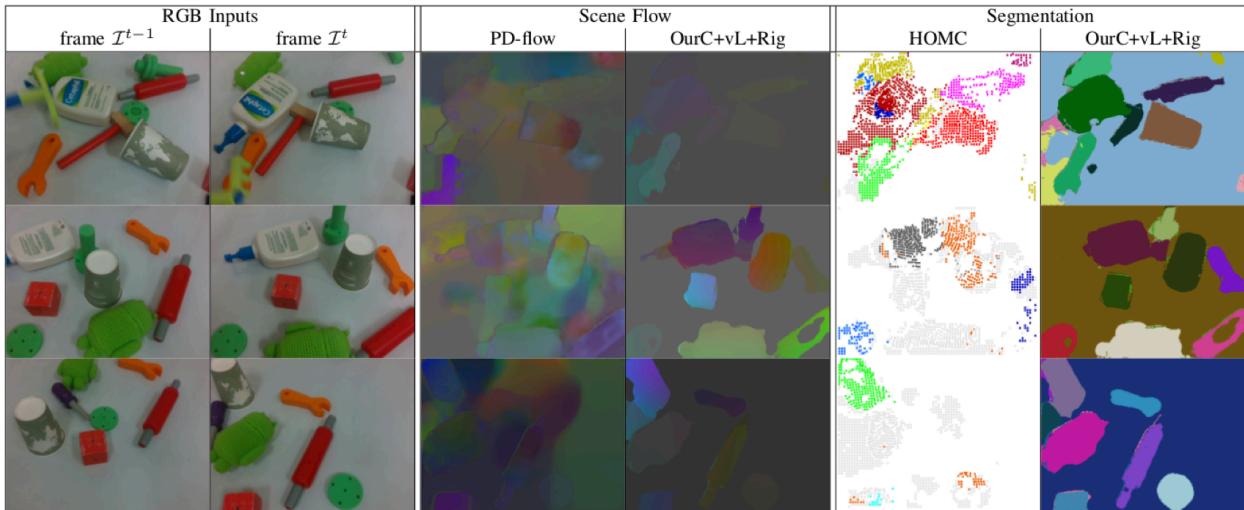


# Motion-based Object Segmentation based on Dense RGB-D Scene Flow



Motion-based Object Segmentation based on Dense RGB-D Scene Flow. Shao et al. RAL + IROS. 2018. Pre-print on arXiv.

# Motion-based Object Segmentation based on Dense RGB-D Scene Flow



Motion-based Object Segmentation based on Dense RGB-D Scene Flow. Shao et al. RAL + IROS. 2018. Pre-print on arXiv.

# CS231

## Introduction to Computer Vision



Next lecture:

Optimal Recursive Estimation