Non-Parametric Few-Shot Learning

CS 330

Logistics

Homework 1 due tonight, Homework 2 out **soon**

Fill out project group form if you haven't already.

Project suggestions & project spreadsheet posted

Plan for Today

Non-Parametric Few-Shot Learning
- Siamese networks, matching networks, prototypical networks
- Case study of few-shot medical image diagnosis

Properties of Meta-Learning Algorithms
- Comparison of approaches

Example Meta-Learning Applications
- Imitation learning, drug discovery, motion prediction, language generation

Goals for by the end of lecture:
- Basics of non-parametric few-shot learning techniques (& how to implement)
- Trade-offs between black-box, optimization-based, and non-parametric meta-learning
- Familiarity with applied formulations of meta-learning

Recap: Black-Box Meta-Learning

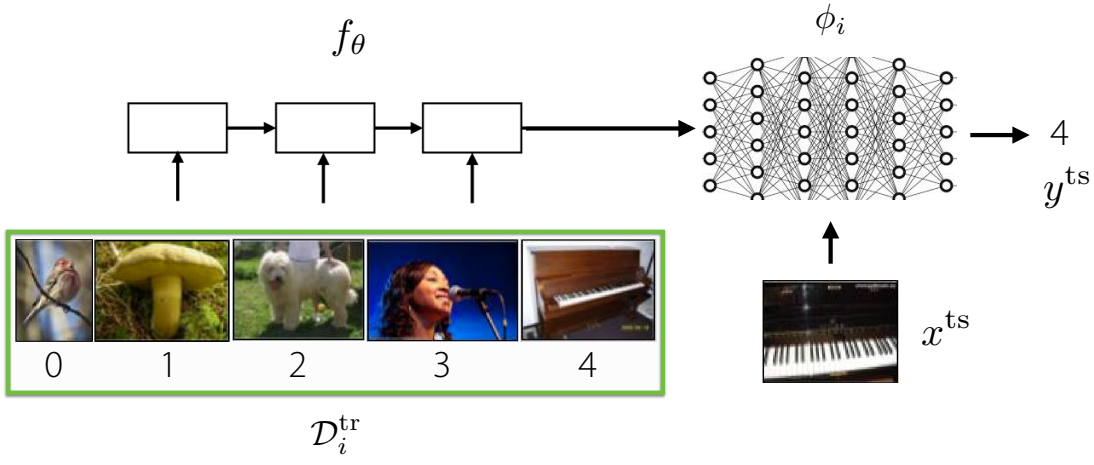

Key idea: parametrize learner as a neural network

+ **expressive** - **challenging optimization** problem

Recap: Optimization-Based Meta-Learning

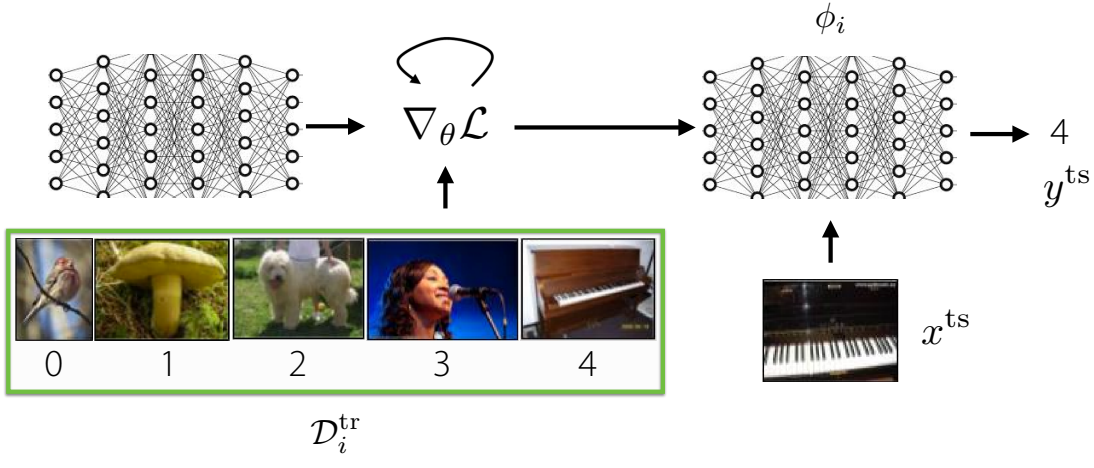

Key idea: embed optimization inside the inner learning process

+ **structure** of **optimization** embedded into meta-learner

- typically requires **second-order optimization**

Today: Can we embed a learning procedure ***without*** a second-order optimization?

So far: Learning parametric models.

There are no parameters involved. data itself are parameters e.g. nearest neighbor

In low data regimes, **non-parametric** methods are simple, work well.

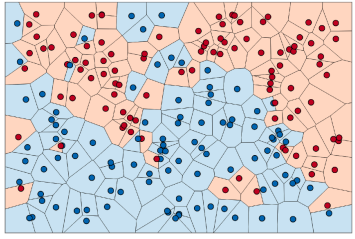

During **meta-test time**: few-shot learning <-> low data regime

During **meta-training**: still want to be parametric

Can we use **parametric meta-learners** that produce effective **non-parametric learners**?

Note: some of these methods precede parametric approaches

Non-parametric methods

Key Idea: Use non-parametric learner.

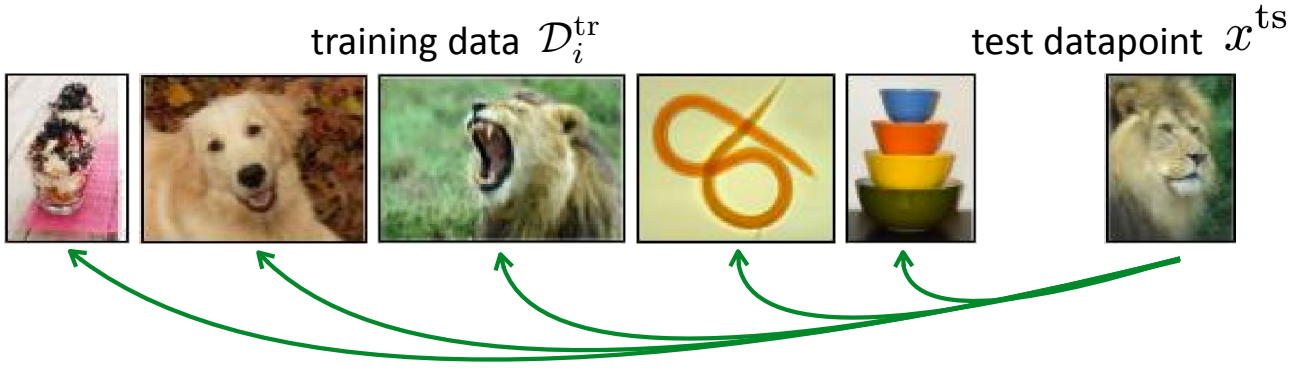

Compare test image with training images

In what space do you compare? With what distance metric?

pixel space, l_2 distance?

In what space do you compare? With what distance metric?

pixel space, l_2 distance?

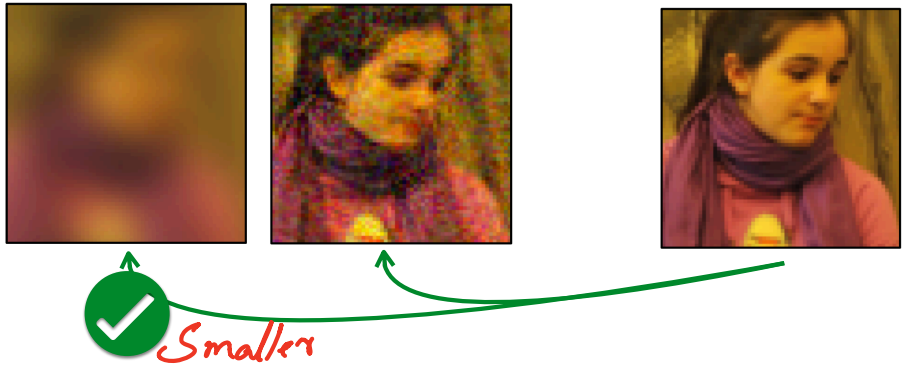

Zhang et al. (arXiv 1801.03924)

Non-parametric methods

Key Idea: Use non-parametric learner.

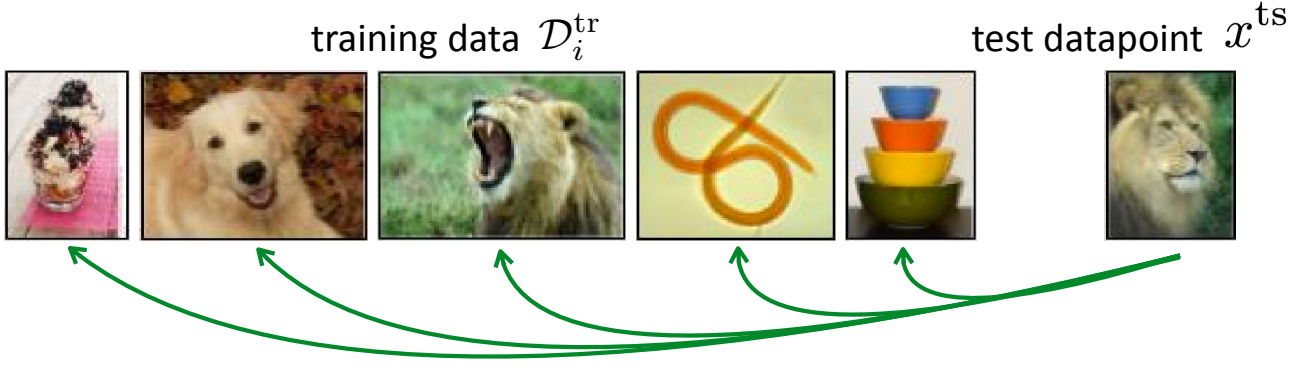

Compare test image with training images

In what space do you compare? With what distance metric?

~~pixel space, l_2 distance?~~

Learn to compare using meta-training data!

Non-parametric methods

Key Idea: Use non-parametric learner.

train Siamese network to predict whether or not two images are the same class

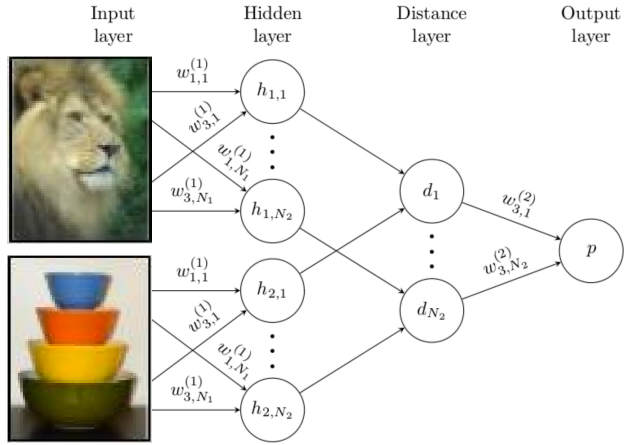

label

0

Koch et al., ICML ‘15

Non-parametric methods

Key Idea: Use non-parametric learner.

train Siamese network to predict whether or not two images are the same class

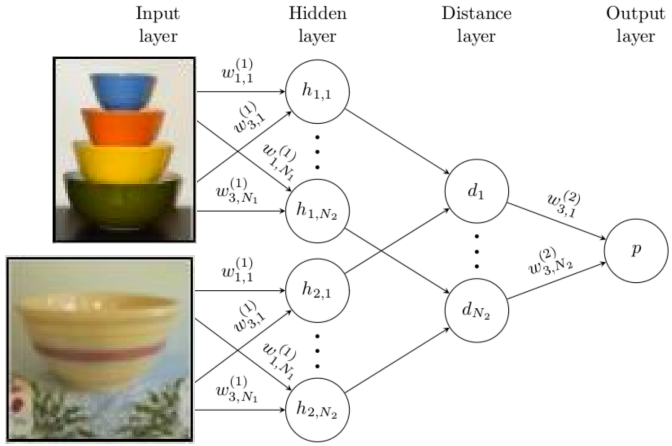

Koch et al., ICML '15

Non-parametric methods

Key Idea: Use non-parametric learner.

train Siamese network to predict whether or not two images are the same class

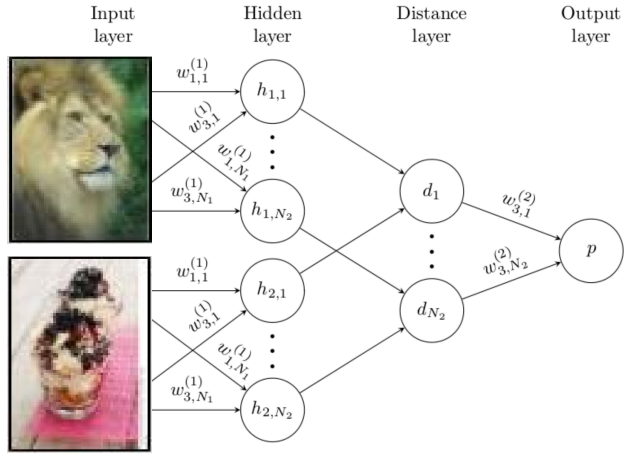

label

0

Koch et al., ICML ‘15

Non-parametric methods

Key Idea: Use non-parametric learner.

train Siamese network to predict whether or not two images are the same class

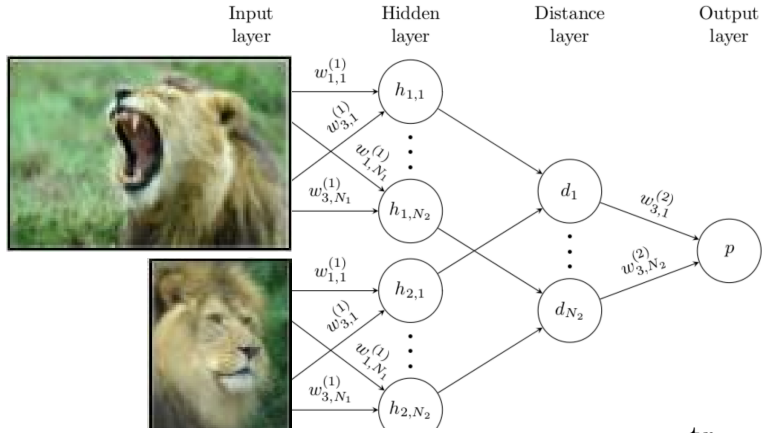

Meta-test time: compare image $\mathbf{x}_{\text{test}}$ to each image in $\mathcal{D}_j^{\text{tr}}$

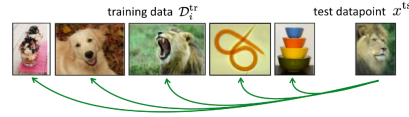

Meta-training: Binary classification

Meta-test: N-way classification

Can we **match** meta-train & meta-test?

Koch et al., ICML '15

Non-parametric methods

ConvLSTM can replace Conv2D+LSTM and is more suitable.

Key Idea: Use non-parametric learner.

Can we **match** meta-train & meta-test?

Nearest neighbors in learned embedding space

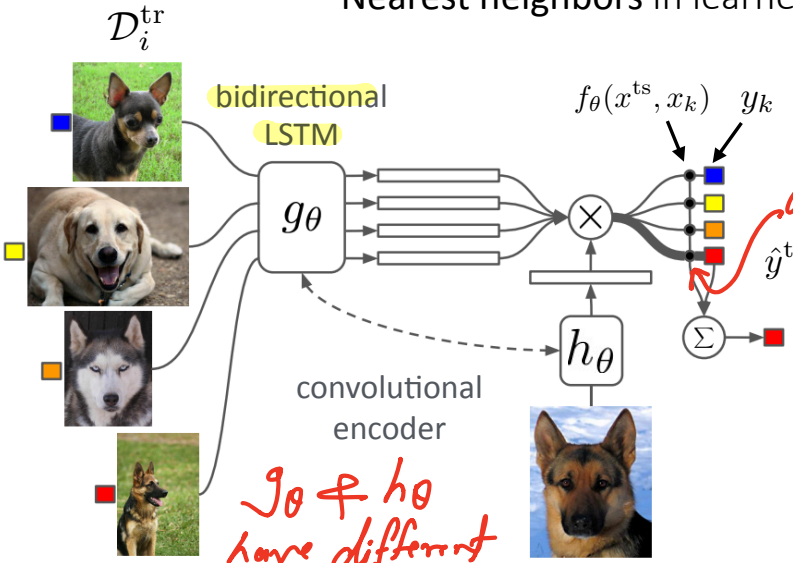

dot product

$$\hat{y}^{\text{ts}} = \sum_{x_k, y_k \in \mathcal{D}^{\text{tr}}} f_\theta(x^{\text{ts}}, x_k) y_k$$

g_θ & h_θ have different architectures

Trained **end-to-end**.

Meta-train & meta-test time **match**.

Vinyals et al. Matching Networks, NeurIPS '16

Non-parametric methods

Key Idea: Use non-parametric learner.

General Algorithm:

~~Black box approach~~ Non-parametric approach (matching networks)

1. Sample task $\mathcal{T}_i$ *(or mini batch of tasks)*
2. Sample disjoint datasets $\mathcal{D}_i^{\text{tr}}, \mathcal{D}_i^{\text{test}}$ from $\mathcal{D}_i$
3. ~~Compute $\phi_i \leftarrow f_\theta(\mathcal{D}_i^{\text{tr}})$~~ Compute $\hat{y}^{\text{ts}} = \sum_{x_k, y_k \in \mathcal{D}^{\text{tr}}} f_\theta(x^{\text{ts}}, x_k) y_k$
4. ~~Update θ using $\nabla_\theta \mathcal{L}(\phi_i, \mathcal{D}_i^{\text{test}})$~~ Update θ using $\nabla_\theta \mathcal{L}(\hat{y}^{\text{ts}}, y^{\text{ts}})$

(Parameters ϕ integrated out, hence **non-parametric**)

What if >1 shot?

Matching networks will perform comparisons **independently**

Can we **aggregate class information** to create a **prototypical embedding**?

What if we had more than 1 example for each class?

e.g. two dogs of same breed.

Non-parametric methods

Key Idea: Use non-parametric learner.

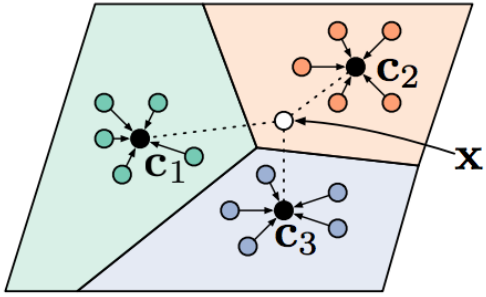

Centroid

$$\mathbf{c}_n = \frac{1}{K} \sum_{(x,y) \in \mathcal{D}_i^{\text{tr}}} \mathbb{1}(y = n) f_\theta(x)$$

$$p_\theta(y = n|x) = \frac{\exp(-d\left(f_\theta(x), \mathbf{c}_n\right))}{\sum_{n'} \exp(d(f_\theta(x), \mathbf{c}_{n'}))}$$

d: Euclidean, or cosine distance

Snell et al. Prototypical Networks, NeurIPS '17

Non-parametric methods

So far: Siamese networks, matching networks, prototypical networks

Embed, then nearest neighbors.

Challenge

What if you need to reason about more complex relationships between datapoints?

Idea: Learn non-linear relation module on embeddings

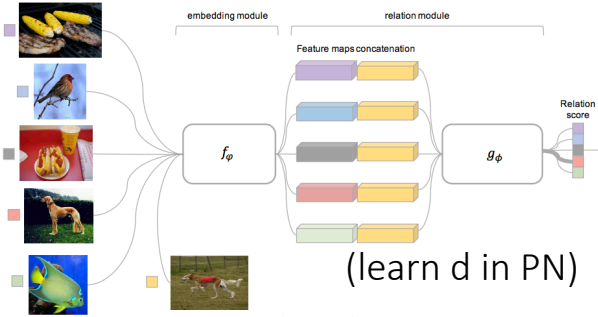

(learn d in PN)

Sung et al. Relation Net

Idea: Learn infinite mixture of prototypes.

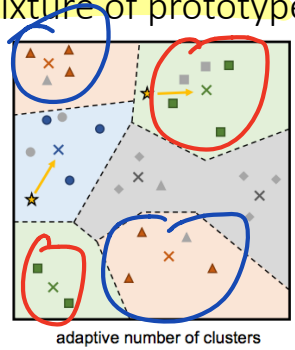

Allen et al. IMP, ICML '19

Idea: Perform message passing on embeddings

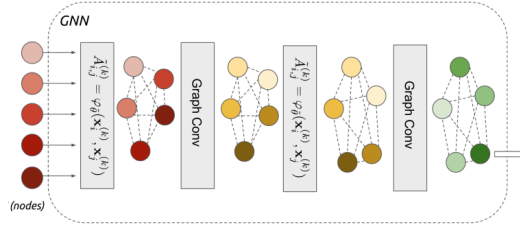

Garcia & Bruna, GNN

Case Study

Prototypical Clustering Networks for Dermatological Image Classification

Viraj Prabhu [*,1] Anitha Kannan[3] Murali Ravuri[3] Manish Chablani[3]
David Sontag[2] Xavier Amatriain[3]
[1]Georgia Tech [2]MIT [3]Curai
virajp@gatech.edu dsontag@mit.edu {anitha, murali, manish, xavier}@curai.com

Machine Learning for Healthcare Conference 2019

NeurIPS 2018 ML4H Workshop

Link: https://arxiv.org/abs/1811.03066

Problem: Few-Shot Learning for Dermatological Disease Diagnosis

Challenges:

- hard to get data
- data is long-tailed
- significant intra-class variability

Goal:

Acquire accurate classifier on all classes

Dermnet dataset

(http://www.dermnet.com/)

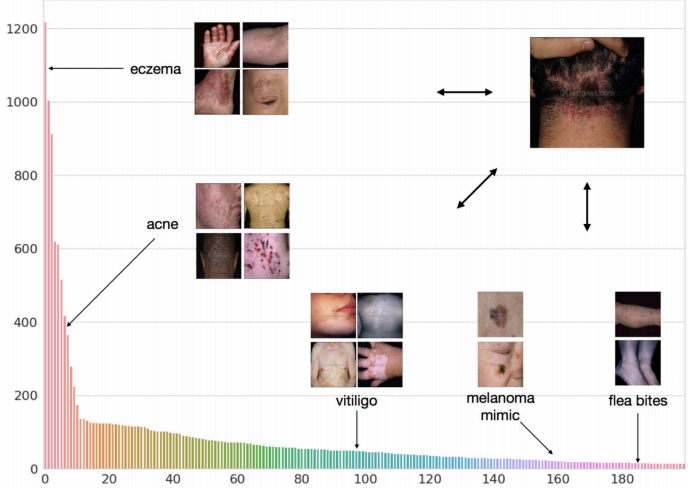

(Top 200 classes only!)

Prototypical Clustering Networks for Few-Shot Classification

Problem formulation:

different image classes = different diseases

150 base classes (classes w/ most data)

50 novel classes

Test on all 200 classes.

Approach: Prototypical Networks +

- learn multiple prototypes per class (to handle intra-class variability)
- incorporate unlabeled support examples via k-means on learned embedding

Note: Unlike black-box & optimization-based meta-learning,
ProtoNets can train for N way classification and test for > N way classification

(Side note if you read the paper: They flipped the standard notation of K and N in the paper)

Evaluation

Compare: PN - standard ProtoNets, trained on 150 base classes, pre-trained on ImageNet

FT_N-*NN - ImageNet pre-training, fine-tuned ResNet on N classes,
*-nearest neighbors in resulting embedding space

FT_{200}-*CE - ImageNet pre-trained, fine-tuned on all 200 classes with balancing

(very strong baseline, accesses more info during training, requires re-training for new classes)

Evaluation Metric: mean class accuracy (mca), i.e. average of per-class accuracies across 200 classes.

	k = 5			k = 10		
Approach	$mca_{base+novel}$	mca_{base}	mca_{novel}	$mca_{base+novel}$	mca_{base}	mca_{novel}
FT_{150}-1NN	46.18 +/- 0.81	55.32 +/- 0.30	18.76 +/- 3.30	49.51 +/- 0.34	54.86 +/- 0.50	33.44 +/- 1.35
FT_{150}-3NN	44.28 +/- 0.32	54.77 +/- 0.47	12.80 +/- 1.50	47.01 +/- 0.56	54.13 +/- 0.43	25.64 +/- 1.51
FT_{200}-1NN	46.52 +/- 0.39	54.17 +/- 0.30	22.50 +/- 0.75	49.92 +/- 0.47	53.80 +/- 0.35	38.27 +/- 1.32
FT_{200}-3NN	44.69 +/- 0.39	52.61 +/- 0.21	20.93 +/- 2.00	47.96 +/- 0.11	52.53 +/- 0.14	34.27 +/- 0.19
FT_{200}-CE	**47.82 +/- 0.46**	**55.75 +/- 0.71**	24.00 +/- 3.22	**51.51 +/- 0.41**	**55.21 +/- 0.26**	40.40 +/- 2.36
PN	43.92 +/- 0.40	48.71 +/- 0.37	29.56 +/- 2.35	44.93 +/- 0.79	47.55 +/- 0.37	37.08 +/- 3.39
PCN (ours)	**47.79 +/- 0.71**	53.70 +/- 0.18	**30.04 +/- 2.77**	**50.92 +/- 0.63**	51.38 +/- 0.34	**49.56 +/- 2.76**

PCN > PN

PCN > FT_N-*NN

PCN ≈ FT_{200}-*CE without requiring re-training

PCN does better on novel classes.

More visualizations and analysis in the paper!

Plan for Today

Non-Parametric Few-Shot Learning
- Siamese networks, matching networks, prototypical networks
- Case study of few-shot medical image diagnosis

Properties of Meta-Learning Algorithms
- Comparison of approaches

Example Meta-Learning Applications
- Imitation learning, drug discovery, motion prediction, language generation

How can we think about how these methods compare?

Black-box vs. Optimization vs. Non-Parametric

***Computation graph* perspective**

Black-box

$$y^{\text{ts}} = f_\theta(\mathcal{D}_i^{\text{tr}}, x^{\text{ts}})$$

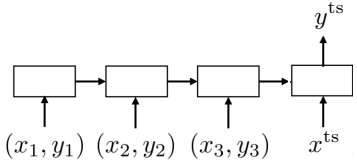

Optimization-based

$$y^{\text{ts}} = f_{\text{MAML}}(\mathcal{D}_i^{\text{tr}}, x^{\text{ts}}) = f_{\phi_i}(x^{\text{ts}})$$

$$\text{where } \phi_i = \theta - \alpha \nabla_\theta \mathcal{L}(\theta, \mathcal{D}_i^{\text{tr}})$$

Non-parametric

$$y^{\text{ts}} = f_{\text{PN}}(\mathcal{D}_i^{\text{tr}}, x^{\text{ts}}) = \text{softmax}(-d\left(f_\theta(x^{\text{ts}}), \mathbf{c}_n\right))$$

$$\text{where } \mathbf{c}_n = \frac{1}{K} \sum_{(x,y) \in \mathcal{D}_i^{\text{tr}}} \mathbb{1}(y = n) f_\theta(x)$$

Note: (again) Can mix & match components of computation graph

1. Both condition on data & run gradient descent.

Jiang et al. CAML '19

2. Gradient descent on relation net embedding.

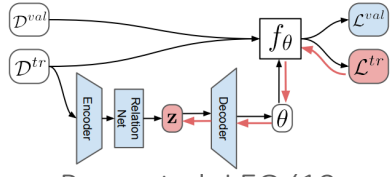

Rusu et al. LEO '19

3. MAML, but initialize last layer as ProtoNet during meta-training

Triantafillou et al. Proto-MAML '19

Black-box vs. Optimization vs. Non-Parametric

***Algorithmic properties* perspective**

Expressive power

the ability for f to represent a range of learning procedures

Why? scalability, applicability to a range of domains

Consistency

learned learning procedure will monotonically improve with more data

Why? reduce reliance on meta-training tasks, good OOD task performance

Recall:

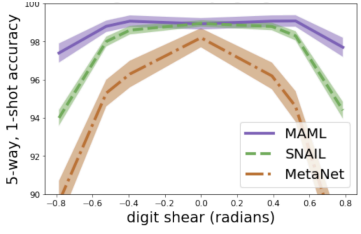

These properties are important for most applications!

Black-box vs. Optimization vs. Non-Parametric

Black-box

+ **complete expressive power**

- **not consistent**

+ easy to combine with **variety of learning problems** (e.g. SL, RL)

- **challenging optimization** (no inductive bias at the initialization)

- often **data-inefficient**

Optimization-based

+ **consistent, reduces** to **GD**

~ **expressive** for **very deep models***

+ **positive inductive bias** at the start of meta-learning

+ handles **varying** & **large K** well

+ **model-agnostic**

- **second-order optimization**

- usually **compute** and **memory** intensive

Non-parametric

+ **expressive** for **most architectures**

~ **consistent** under **certain conditions**

+ entirely **feedforward**

+ **computationally fast** & **easy to optimize**

- **harder to generalize** to **varying K**

- hard to scale to **very large K**

- so far, **limited to classification**

Generally, well-tuned versions of each perform **comparably** on existing few-shot benchmarks!

(likely says more about the benchmarks than the methods)

Which method to use depends on your **use-case**.

for classification, Non-parametric is the easiest to get up and running.

for supervised learning settings

Black-box vs. Optimization vs. Non-Parametric

***Algorithmic properties* perspective**

Expressive power

the ability for f to represent a range of learning procedures

Why? scalability, applicability to a range of domains

Consistency

learned learning procedure will monotonically improve with more data

Why? reduce reliance on meta-training tasks, good OOD task performance

Uncertainty awareness

ability to reason about ambiguity during learning

Why? active learning, calibrated uncertainty, RL, principled Bayesian approaches

We'll discuss this next Weds!

Plan for Today

Non-Parametric Few-Shot Learning
- Siamese networks, matching networks, prototypical networks
- Case study of few-shot medical image diagnosis

Properties of Meta-Learning Algorithms
- Comparison of approaches

Example Meta-Learning Applications
- Imitation learning, drug discovery, motion prediction, language generation

Application: One-Shot Imitation Learning

(Yu*, Finn* et al. One-Shot Imitation from Observing Humans. RSS 2018)

Tasks:

manipulating different objects

$\mathcal{D}_i^{\text{tr}}$: video of a human

$\mathcal{D}_i^{\text{ts}}$: teleoperated demonstration

Model: optimization-based

MAML with ***learned*** inner loss

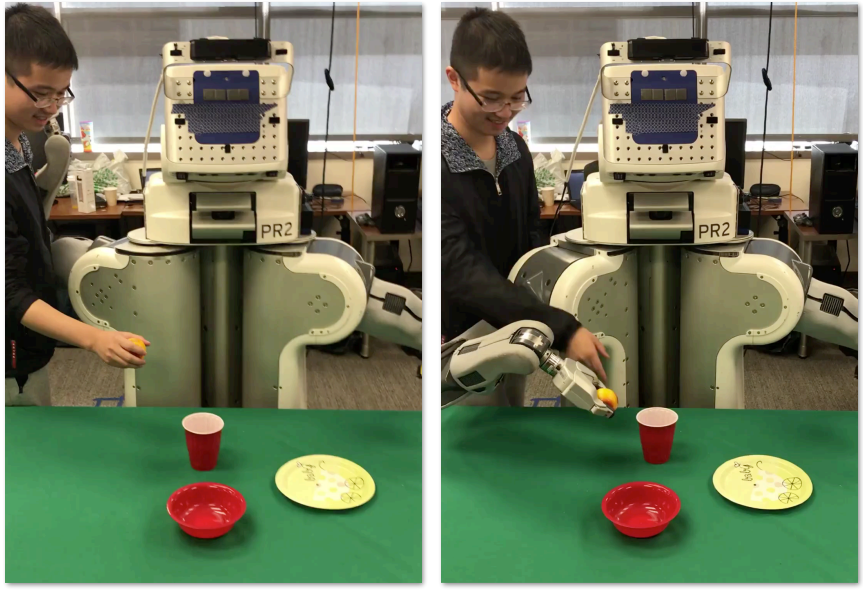

Application: Low-Resource Molecular Property Prediction

(Nguyen et al. Meta-Learning GNN Initializations for Low-Resource Molecular Property Prediction. 2020)

[potentially useful for low-resource drug discovery problems]

Tasks:

Predicting properties & activities of different molecules

$\mathcal{D}_i^{\text{tr}}$, $\mathcal{D}_i^{\text{ts}}$: different instances

Model: optimization-based

MAML, first-order MAML, ANIL

Gated graph neural net base model

Base Model

ChEMBL ID	k-NN	Finetune-All	Finetune-Top	FO-MAML	ANIL	MAML
2363236	0.316 ± 0.007	0.328 ± 0.028	0.329 ± 0.023	$\mathbf{0.337 \pm 0.019}$	0.325 ± 0.008	0.332 ± 0.013
1614469	0.438 ± 0.023	0.470 ± 0.034	0.490 ± 0.033	0.489 ± 0.019	0.446 ± 0.044	$\mathbf{0.507 \pm 0.030}$
2363146	0.559 ± 0.026	0.626 ± 0.037	$\mathbf{0.653 \pm 0.029}$	0.555 ± 0.017	0.506 ± 0.034	0.595 ± 0.051
2363366	0.511 ± 0.050	0.567 ± 0.039	0.551 ± 0.048	0.546 ± 0.037	$\mathbf{0.570 \pm 0.031}$	$\mathbf{0.598 \pm 0.041}$
2363553	$\mathbf{0.739 \pm 0.007}$	0.724 ± 0.015	0.737 ± 0.023	0.694 ± 0.011	0.686 ± 0.020	0.691 ± 0.013
1963818	0.607 ± 0.041	0.708 ± 0.036	0.595 ± 0.142	0.677 ± 0.026	0.692 ± 0.081	$\mathbf{0.745 \pm 0.048}$
1963945	0.805 ± 0.031	$\mathbf{0.848 \pm 0.034}$	0.835 ± 0.036	0.779 ± 0.039	0.753 ± 0.033	0.836 ± 0.023
1614423	0.503 ± 0.044	0.628 ± 0.058	0.642 ± 0.063	0.760 ± 0.024	0.730 ± 0.077	$\mathbf{0.837 \pm 0.036}^*$
2114825	0.679 ± 0.027	0.739 ± 0.050	0.732 ± 0.051	0.837 ± 0.042	0.759 ± 0.078	$\mathbf{0.885 \pm 0.014}^*$
1964116	0.709 ± 0.042	0.758 ± 0.044	0.769 ± 0.048	0.895 ± 0.023	0.903 ± 0.016	$\mathbf{0.912 \pm 0.013}$
2155446	0.471 ± 0.008	0.473 ± 0.017	0.476 ± 0.013	0.497 ± 0.024	0.478 ± 0.020	$\mathbf{0.500 \pm 0.017}$
1909204	0.538 ± 0.023	0.589 ± 0.031	0.577 ± 0.039	0.592 ± 0.043	0.547 ± 0.029	$\mathbf{0.601 \pm 0.027}$
1909213	0.694 ± 0.009	0.742 ± 0.015	$\mathbf{0.759 \pm 0.012}$	0.698 ± 0.024	0.694 ± 0.025	0.729 ± 0.013
3111197	0.617 ± 0.028	0.663 ± 0.066	0.673 ± 0.071	0.636 ± 0.036	0.737 ± 0.035	$\mathbf{0.746 \pm 0.045}$
3215171	0.480 ± 0.042	0.552 ± 0.043	0.551 ± 0.045	0.729 ± 0.031	0.700 ± 0.050	$\mathbf{0.764 \pm 0.019}$
3215034	0.474 ± 0.072	0.540 ± 0.156	0.455 ± 0.189	$\mathbf{0.819 \pm 0.048}$	0.681 ± 0.042	0.805 ± 0.046
1909103	0.881 ± 0.026	$\mathbf{0.936 \pm 0.013}$	0.921 ± 0.020	0.877 ± 0.046	0.730 ± 0.055	0.900 ± 0.032
3215092	0.696 ± 0.038	0.777 ± 0.039	0.791 ± 0.042	0.877 ± 0.028	0.834 ± 0.026	$\mathbf{0.907 \pm 0.017}$
1738253	0.710 ± 0.048	0.860 ± 0.029	0.861 ± 0.025	0.885 ± 0.033	0.758 ± 0.111	$\mathbf{0.908 \pm 0.011}$
1614549	0.710 ± 0.035	0.850 ± 0.041	0.860 ± 0.051	0.930 ± 0.022	0.860 ± 0.034	$\mathbf{0.947 \pm 0.014}$
Avg. Rank	5.4	3.5	3.5	3.1	4.0	**1.7**

Application: Few-Shot Human Motion Prediction

(Gui et al. Few-Shot Human Motion Prediction via Meta-Learning. ECCV 2018)

[potentially useful for human-robot interaction, autonomous driving]

Tasks:

Different human users & motions

$\mathcal{D}_i^{\mathrm{tr}}$: past K time steps of motion

$\mathcal{D}_i^{\mathrm{ts}}$: future second(s) of motion

Model:

optimization-based/black-box hybrid

MAML with additional learned update rule

Recurrent neural net base model

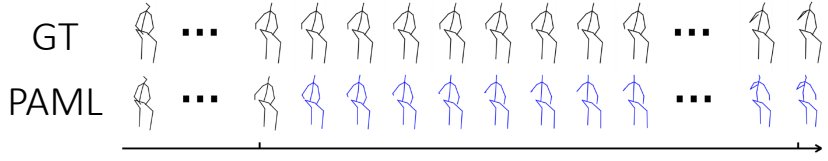

		Walking						Eating					
milliseconds		80	160	320	400	560	1000	80	160	320	400	560	1000
residual sup. [32] w/ (Baselines)	$\text{Scratch}_{\text{spec}}$	1.90	1.95	2.16	2.18	1.99	2.00	2.33	2.31	2.30	2.30	2.31	2.34
	$\text{Scratch}_{\text{agn}}$	1.78	1.89	2.20	2.23	2.02	2.05	2.27	2.16	2.18	2.27	2.25	2.31
	$\text{Transfer}_{\text{ots}}$	0.60	0.75	0.88	0.93	1.03	1.26	0.57	0.70	0.91	1.04	1.19	1.58
	Multi-task	0.57	0.71	0.79	0.85	0.96	1.12	0.59	0.68	0.83	0.93	1.12	1.33
	$\text{Transfer}_{\text{ft}}$	0.44	0.55	0.85	0.95	0.74	1.03	0.61	0.65	0.74	0.78	0.86	1.19
Meta-learning (Ours)	PAML	**0.35**	**0.47**	**0.70**	**0.82**	**0.80**	**0.83**	**0.36**	**0.52**	**0.65**	**0.70**	**0.71**	**0.79**
		Smoking						Discussion					
milliseconds		80	160	320	400	560	1000	80	160	320	400	560	1000
residual sup. [32] w/ (Baselines)	$\text{Scratch}_{\text{spec}}$	2.88	2.86	2.85	2.83	2.80	2.99	3.01	3.13	3.12	2.95	2.62	2.99
	$\text{Scratch}_{\text{agn}}$	2.53	2.61	2.67	2.65	2.71	2.73	2.77	2.79	2.82	2.73	2.82	2.76
	$\text{Transfer}_{\text{ots}}$	0.70	0.84	1.18	1.23	1.38	2.02	0.58	0.86	1.12	1.18	1.54	2.02
	Multi-task	0.71	0.79	1.09	1.20	1.25	1.23	0.53	0.82	1.02	1.17	1.33	1.97
	$\text{Transfer}_{\text{ft}}$	0.87	1.02	1.25	1.30	1.45	2.06	0.57	0.82	1.11	1.11	1.37	2.08
Meta-learning (Ours)	PAML	**0.39**	**0.66**	**0.81**	**1.01**	**1.03**	**1.01**	**0.41**	**0.71**	**1.01**	**1.02**	**1.09**	**1.12**

mean angle error w.r.t. prediction horizon

Application: Language Modeling

(Brown*, Mann*, Ryder*, Subbiah* et al. Language Models are Few-Shot Learners. 2020)

Tasks:

spelling correction
simple math problems
language translation
a variety of others

All represented as language generation problems

$\mathcal{D}_i^{\text{tr}}$: sequence of characters

$\mathcal{D}_i^{\text{ts}}$: following sequence of characters

Model: black-box meta-learner

giant "Transformer" model

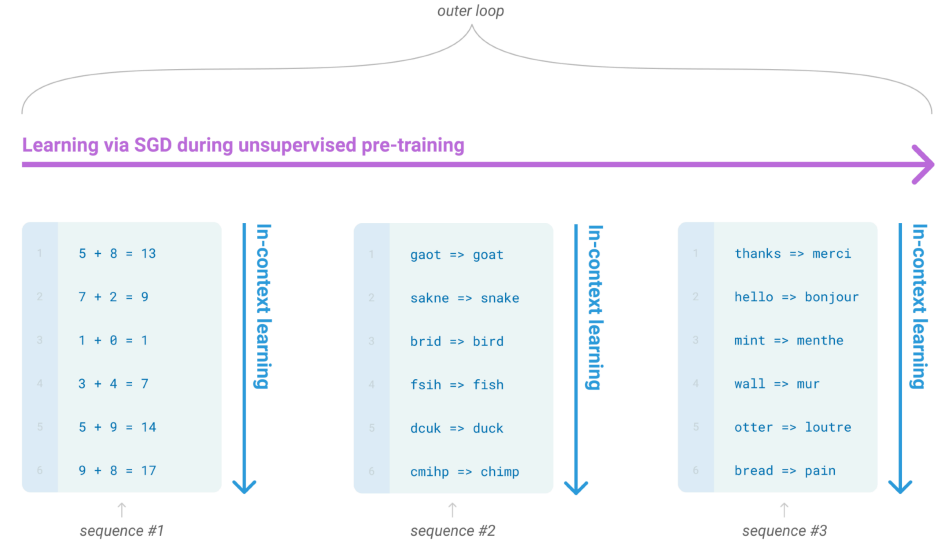

Some Results

One-shot learning from dictionary definitions:

To "screeg" something is to swing a sword at it. An example of a sentence that uses the word screeg is:
We screeghed at each other for several minutes and then we went outside and ate ice cream.

Few-shot language editing:

Poor English input: I eated the purple berries.
Good English output: I ate the purple berries.
Poor English input: Thank you for picking me as your designer. I'd appreciate it.
Good English output: Thank you for choosing me as your designer. I appreciate it.
Poor English input: The mentioned changes have done. or I did the alteration that you requested. or I changed things you wanted and did the modifications.
Good English output: The requested changes have been made. or I made the alteration that you requested. or I changed things you wanted and made the modifications.
Poor English input: I'd be more than happy to work with you in another project.
Good English output: I'd be more than happy to work with you on another project.

Poor English input: Please provide me with a short brief of the design you're looking for and that'd be nice if you could share some examples or project you did before.
Good English output: Please provide me with a brief description of the design you're looking for and that would be nice if you could share some examples or projects you have done before.

Non-few-shot learning tasks:

Title: United Methodists Agree to Historic Split
Subtitle: Those who oppose gay marriage will form their own denomination
Article: **After two days of intense debate, the United Methodist Church has agreed to a historic split - one that is expected to end in the creation of a new denomination, one that will be "theologically and socially conservative," according to The Washington Post. The majority of delegates attending the church's annual General Conference in May voted to strengthen a ban on the ordination of LGBTQ clergy and to write new rules that will "discipline" clergy who officiate at same-sex weddings. But those who opposed these measures have a new plan: They say they will form a**

Plan for Today

Non-Parametric Few-Shot Learning
- Siamese networks, matching networks, prototypical networks
- Case study of few-shot medical image diagnosis

Properties of Meta-Learning Algorithms
- Comparison of approaches

Example Meta-Learning Applications
- Imitation learning, drug discovery, motion prediction, language generation

Goals for by the end of lecture:
- Basics of non-parametric few-shot learning techniques (& how to implement)
- Trade-offs between black-box, optimization-based, and non-parametric meta-learning
- Familiarity with applied formulations of meta-learning

Reminders

Homework 1 due tonight, Homework 2 out **soon**

Fill out project group form if you haven't already.

Project suggestions & project spreadsheet posted