

Reinforcement Learning: A Primer, Multi-Task, Goal-Conditioned

CS 330

Introduction

Some background:

- Not a native English speaker so **please please** let me know if you don't understand something
- I like robots 😊
- Studied classical robotics first
- Got fascinated by deep RL in the middle of my PhD after a talk by Sergey Levine
- Research Scientist at Robotics @ Google




Karol
Hausman



Why Reinforcement Learning?

Isolated action that doesn't affect the future?



Why Reinforcement Learning?

Isolated action that doesn't affect the future?

Supervised learning?

In RL, the actions effect future behaviour.

Common applications



robotics



language & dialog



autonomous driving



business operations



finance

(most deployed ML systems)

+ a key aspect of intelligence

The Plan

Multi-task reinforcement learning problem

Policy gradients & their multi-task counterparts

Q-learning

<— should be review

Multi-task Q-learning

The Plan

Multi-task reinforcement learning problem

Policy gradients & their multi-task counterparts

Q-learning


<— should be review

Multi-task Q-learning

Terminology & notation



\mathbf{o}_t



$\pi_\theta(\mathbf{a} | \mathbf{o}_t)$

$\pi_\theta(a_t | o_t)$



\mathbf{a}_t

1. run away
2. ignore
3. pet

\mathbf{s}_t – state

\mathbf{o}_t – observation

\mathbf{a}_t – action

$\pi_\theta(\mathbf{a}_t | \mathbf{o}_t)$ – policy

$\pi_\theta(\mathbf{a}_t | \mathbf{s}_t)$ – policy (fully observed)




\mathbf{o}_t – observation




\mathbf{s}_t – state

Terminology & notation




Imitation Learning



Images: Bojarski et al. '16, NVIDIA

Slide adapted from Sergey Levine

Imitation Learning




Imitation Learning vs Reinforcement Learning?



Images: Bojarski et al. '16, NVIDIA

Slide adapted from Sergey Levine

Reward functions



$$\pi_{\theta}(\mathbf{a}_t | \mathbf{o}_t)$$



which action is better or worse?

$r(\mathbf{s}, \mathbf{a})$: reward function

tells us which states and actions are better





high reward



low reward

\mathbf{s} , \mathbf{a} , $r(\mathbf{s}, \mathbf{a})$, and $p(\mathbf{s}'|\mathbf{s}, \mathbf{a})$ define
Markov decision process


The goal of reinforcement learning





Partial observability

Fully observable?

- Simulated robot performing a reaching task given the goal position and positions and velocities of all of its joints
- Indiscriminate robotic grasping from a bin given an overhead image
- A robot sorting trash given a camera image



The goal of reinforcement learning



$$\theta^* = \arg \max_{\theta} E_{(\mathbf{s}, \mathbf{a}) \sim p_\theta(\mathbf{s}, \mathbf{a})}[r(\mathbf{s}, \mathbf{a})]$$

infinite horizon case

$$\theta^* = \arg \max_{\theta} \sum_{t=1}^T E_{(\mathbf{s}_t, \mathbf{a}_t) \sim p_\theta(\mathbf{s}_t, \mathbf{a}_t)}[r(\mathbf{s}_t, \mathbf{a}_t)]$$

finite horizon case

What is a reinforcement learning **task**?

Recall: supervised learning

data generating distributions, loss

A task: $\mathcal{T}_i \triangleq \{p_i(\mathbf{x}), p_i(\mathbf{y}|\mathbf{x}), \mathcal{L}_i\}$

Reinforcement learning

action space

dynamics

A task: $\mathcal{T}_i \triangleq \{\mathcal{S}_i, \mathcal{A}_i, p_i(\mathbf{s}_1), p_i(\mathbf{s}'|\mathbf{s}, \mathbf{a}), r_i(\mathbf{s}, \mathbf{a})\}$



state
space

initial state
distribution



reward


a Markov decision process

much more than the semantic meaning of task!


Examples Task Distributions

A task: $\mathcal{T}_i \triangleq \{\mathcal{S}_i, \mathcal{A}_i, p_i(\mathbf{s}_1), p_i(\mathbf{s}'|\mathbf{s}, \mathbf{a}), r_i(\mathbf{s}, \mathbf{a})\}$

Character animation:
across maneuvers
 $r_i(\mathbf{s}, \mathbf{a})$ vary



across garments &
initial states
 $p_i(\mathbf{s}_1), p_i(\mathbf{s}'|\mathbf{s}, \mathbf{a})$ vary




Multi-robot RL:



$\mathcal{S}_i, \mathcal{A}_i, p_i(\mathbf{s}_1), p_i(\mathbf{s}'|\mathbf{s}, \mathbf{a})$ vary

What is a reinforcement learning **task**?



An alternative view:


A task identifier is part of the state: $\mathbf{s} = (\bar{\mathbf{s}}, \mathbf{z}_i)$

\swarrow
original state

$$\mathcal{T}_i \triangleq \{\mathcal{S}_i, \mathcal{A}_i, p_i(\mathbf{s}_1), p(\mathbf{s}'|\mathbf{s}, \mathbf{a}), r(\mathbf{s}, \mathbf{a})\} \longrightarrow \{\mathcal{T}_i\} = \left\{ \cup \mathcal{S}_i, \cup \mathcal{A}_i, \frac{1}{N} \sum_i p_i(\mathbf{s}_1), p(\mathbf{s}'|\mathbf{s}, \mathbf{a}), r(\mathbf{s}, \mathbf{a}) \right\}$$

It can be cast as a **standard Markov decision process!**

The goal of multi-task reinforcement learning



Multi-task RL

The same as before, except:

a task identifier is part of the state: $s = (\bar{s}, z_i)$

e.g. one-hot task ID

language description

desired goal state, $z_i = s_g \leftarrow$ “goal-conditioned RL”

If it's still a standard **Markov decision process**,

then, why not apply standard **RL algorithms**?

You can!

You can often do better.

What is the reward?

The same as before

Or, for **goal-conditioned RL**:

$$r(s) = r(\bar{s}, s_g) = -d(\bar{s}, s_g)$$

Distance function d examples:

- Euclidean ℓ_2
- sparse 0/1

The Plan


Multi-task reinforcement learning problem

Policy gradients & their multi-task counterparts

Q-learning

Multi-task Q-learning

The anatomy of a reinforcement learning algorithm



This lecture: focus on model-free RL methods (policy gradient, Q-learning)

10/19: focus on model-based RL methods

On-policy

vs

Off-policy

- Data comes from the current policy
- Compatible with all RL algorithms
- Can't reuse data from previous policies

- Data comes from any policy
- Works with specific RL algorithms
- Much more sample efficient, can re-use old data


Evaluating the objective *(Policy-gradient)*

$$\theta^* = \arg \max_{\theta} E_{\tau \sim p_{\theta}(\tau)} \left[\sum_t r(\mathbf{s}_t, \mathbf{a}_t) \right]$$

$J(\theta)$

$$J(\theta) = E_{\tau \sim p_{\theta}(\tau)} \left[\sum_t r(\mathbf{s}_t, \mathbf{a}_t) \right] \approx \frac{1}{N} \sum_i \sum_t r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t})$$

sum over samples from π_{θ}



Direct policy differentiation

$$\theta^* = \arg \max_{\theta} E_{\tau \sim p_{\theta}(\tau)} \left[\sum_t r(\mathbf{s}_t, \mathbf{a}_t) \right]$$

$\underbrace{\qquad\qquad\qquad}_{J(\theta)}$

a convenient identity

$$\underline{\pi_{\theta}(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau)} = \pi_{\theta}(\tau) \frac{\nabla_{\theta} \pi_{\theta}(\tau)}{\pi_{\theta}(\tau)} = \underline{\nabla_{\theta} \pi_{\theta}(\tau)}$$

$$J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)} [r(\tau)] = \int \pi_{\theta}(\tau) r(\tau) d\tau$$

$\sum_{t=1}^T \underbrace{r(\mathbf{s}_t, \mathbf{a}_t)}_{}$

$$\nabla_x \log f(x) = \frac{1}{f(x)} \nabla_x f(x)$$

⇒ $(\nabla_x f(x)) = f'(x) \nabla_x \log f(x)$

$$\nabla_{\theta} J(\theta) = \int \underline{\nabla_{\theta} \pi_{\theta}(\tau)} r(\tau) d\tau = \int \underline{\pi_{\theta}(\tau) \nabla_{\theta} \log \pi_{\theta}(\tau)} r(\tau) d\tau = E_{\tau \sim \pi_{\theta}(\tau)} [\nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau)]$$

Direct policy differentiation

$$\theta^* = \arg \max_{\theta} J(\theta)$$

$$J(\theta) = E_{\tau \sim \pi_\theta(\tau)}[r(\tau)]$$

$$\nabla_{\theta} J(\theta) = E_{\tau \sim \pi_\theta(\tau)}[\nabla_{\theta} \log \pi_{\theta}(\tau) r(\tau)]$$

$$\pi_{\theta}(\tau) = \underbrace{\pi_{\theta}(s_1, a_1, \dots, s_T, a_T)}_{\log \text{ of both sides}} = p(s_1) \prod_{t=1}^T \pi_{\theta}(a_t | s_t) p(s_{t+1} | s_t, a_t)$$

$$\log \pi_{\theta}(\tau) = \log p(s_1) + \sum_{t=1}^T \log \pi_{\theta}(a_t | s_t) + \log p(s_{t+1} | s_t, a_t)$$

$$\nabla_{\theta} \left[\cancel{\log p(s_1)} + \sum_{t=1}^T \log \pi_{\theta}(a_t | s_t) + \cancel{\log p(s_{t+1} | s_t, a_t)} \right]$$

does not depend on θ

$$\boxed{\nabla_{\theta} J(\theta) = E_{\tau \sim \pi_\theta(\tau)} \left[\left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right) \left(\sum_{t=1}^T r(s_t, a_t) \right) \right]}$$

Evaluating the policy gradient

recall: $J(\theta) = E_{\tau \sim p_\theta(\tau)} \left[\sum_t r(\mathbf{s}_t, \mathbf{a}_t) \right] \approx \frac{1}{N} \sum_i \sum_t r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t})$



$$\nabla_\theta J(\theta) = E_{\tau \sim \pi_\theta(\tau)} \left[\left(\sum_{t=1}^T \nabla_\theta \log \pi_\theta(\mathbf{a}_t | \mathbf{s}_t) \right) \left(\sum_{t=1}^T r(\mathbf{s}_t, \mathbf{a}_t) \right) \right]$$

$$\nabla_\theta J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_\theta \log \pi_\theta(\mathbf{a}_{i,t} | \mathbf{s}_{i,t}) \right) \left(\sum_{t=1}^T r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) \right)$$

$$\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta) \rightarrow \text{gradient ascent}$$

REINFORCE algorithm:

- 1. sample $\{\tau^i\}$ from $\pi_\theta(\mathbf{a}_t | \mathbf{s}_t)$ (run the policy)
- 2. $\nabla_\theta J(\theta) \approx \sum_i (\sum_t \nabla_\theta \log \pi_\theta(\mathbf{a}_t^i | \mathbf{s}_t^i)) (\sum_t r(\mathbf{s}_t^i, \mathbf{a}_t^i))$
- 3. $\theta \leftarrow \theta + \alpha \nabla_\theta J(\theta)$




Comparison to maximum likelihood

$$\text{policy gradient: } \nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_{i,t} | \mathbf{s}_{i,t}) \right) \left(\sum_{t=1}^T r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) \right)$$

Multi-task learning algorithms
can readily be applied!

$$\text{maximum likelihood: } \nabla_{\theta} J_{\text{ML}}(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_{i,t} | \mathbf{s}_{i,t}) \right)$$



\mathbf{s}_t

$\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)$



\mathbf{a}_t



\mathbf{s}_t
 \mathbf{a}_t

training
data

supervised
learning

$\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)$

What did we just do?

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_{i,t} | \mathbf{s}_{i,t}) \right) \left(\sum_{t=1}^T r(\mathbf{s}_{i,t}, \mathbf{a}_{i,t}) \right)$$

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{i=1}^N \underbrace{\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(\tau_i)}_{r(\tau_i)} r(\tau_i)$$

maximum likelihood: $\nabla_{\theta} J_{\text{ML}}(\theta) \approx \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \log \pi_{\theta}(\tau_i)$


good stuff is made more likely

bad stuff is made less likely

simply formalizes the notion of “trial and error”!

REINFORCE algorithm:

- 1. sample $\{\tau^i\}$ from $\pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t)$ (run it on the robot)
- 2. $\nabla_{\theta} J(\theta) \approx \sum_i (\sum_t \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_t^i | \mathbf{s}_t^i)) (\sum_t r(\mathbf{s}_t^i, \mathbf{a}_t^i))$
- 3. $\theta \leftarrow \theta + \alpha \nabla_{\theta} J(\theta)$



Policy Gradients

$$\text{policy gradient: } \nabla_{\theta} J(\theta) = E_{\tau \sim \pi_{\theta}(\tau)} \left[\left(\sum_{t=1}^T \nabla_{\theta} \log \pi_{\theta}(\mathbf{a}_t | \mathbf{s}_t) \right) \left(\sum_{t=1}^T r(\mathbf{s}_t, \mathbf{a}_t) \right) \right]$$

Pros:

- + Simple
- + Easy to combine with existing multi-task & meta-learning algorithms

Cons:

- Produces a **high-variance gradient**
 - Can be mitigated with baselines (used by all algorithms in practice), trust regions
- Requires **on-policy** data
 - Cannot reuse existing experience to estimate the gradient!
 - Importance weights can help, but also high variance

The Plan

Multi-task reinforcement learning problem

Policy gradients & their multi-task/meta counterparts

Q-learning

Multi-task Q-learning

Value-Based RL: Definitions

Value function: $V^\pi(\mathbf{s}_t) = \sum_{t'=t}^T \mathbb{E}_\pi[r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) \mid \mathbf{s}_t]$ total reward starting from \mathbf{s} and following π
"how good is a state"

Q function: $Q^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^T \mathbb{E}_\pi[r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) \mid \mathbf{s}_t, \mathbf{a}_t]$ total reward starting from \mathbf{s} , taking \mathbf{a} ,
and then following π
"how good is a state-action pair"

They're related: $V^\pi(\mathbf{s}_t) = \mathbb{E}_{\mathbf{a}_t \sim \pi(\cdot | \mathbf{s}_t)}[Q^\pi(\mathbf{s}_t, \mathbf{a}_t)]$

If you know Q^π , you can use it to **improve** π .

Set $\pi(\mathbf{a} | \mathbf{s}) \leftarrow 1$ for $\mathbf{a} = \operatorname{argmax}_{\bar{\mathbf{a}}} Q^\pi(\mathbf{s}, \bar{\mathbf{a}})$ New policy is at least as good as old policy.

Value-Based RL: Definitions

Value function: $V^\pi(\mathbf{s}_t) = \sum_{t'=t}^T \mathbb{E}_\pi[r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) \mid \mathbf{s}_t]$ total reward starting from \mathbf{s} and following π
"how good is a state"

Q function: $Q^\pi(\mathbf{s}_t, \mathbf{a}_t) = \sum_{t'=t}^T \mathbb{E}_\pi[r(\mathbf{s}_{t'}, \mathbf{a}_{t'}) \mid \mathbf{s}_t, \mathbf{a}_t]$ total reward starting from \mathbf{s} , taking \mathbf{a} ,
and then following π
"how good is a state-action pair"

For the optimal policy π^* : $Q^*(\mathbf{s}_t, \mathbf{a}_t) = \mathbb{E}_{\mathbf{s}' \sim p(\cdot | \mathbf{s}, \mathbf{a})} [r(\mathbf{s}, \mathbf{a}) + \gamma \max_{\mathbf{a}'} Q^*(\mathbf{s}', \mathbf{a}')]$

Bellman equation

Value-Based RL


Value function: $V^\pi(s_t) = ?$

Q function: $Q^\pi(s_t, a_t) = ?$

Q* function: $Q^*(s_t, a_t) = ?$

Value* function: $V^*(s_t) = ?$

Reward = 1 if I can play it in a month, 0 otherwise



if,
Current $\pi(a_t | s) = 1$
 $V^\pi(s_t) = 1$

IMPROVISATION TEST EXAMPLES AND IDEAS FOR ROCKSCHOOL GRADE 1 DRUMS EXAM
Written by Theo Lawrence / TL Music Lessons

J=76

Exercise 1: Rock

Exercise 2: Rock

Exercise 3: Rock

Exercise 4: Rock

Exercise 5: Funk Rock

Exercise 6: Rock


Exercise 7: Blues

Exercise 8: Blues

$\delta_{\text{old}} \pi(a_t | s) \leftarrow \pi(a_t^* | s)$ for $a = \arg \max_{\bar{a}} Q^\pi(s, \bar{a})$ New policy is at least as good as old policy.

Fitted Q-iteration Algorithm

full fitted Q-iteration algorithm:

- 
1. collect dataset $\{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)\}$ using some policy
 2. set $\mathbf{y}_i \leftarrow r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \max_{\mathbf{a}'_i} Q_\phi(\mathbf{s}'_i, \mathbf{a}'_i)$
 3. set $\phi \leftarrow \arg \min_\phi \frac{1}{2} \sum_i \|Q_\phi(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{y}_i\|^2$


Algorithm hyperparameters

dataset size N , collection policy

$K \times$

iterations K

gradient steps S



Result: get a policy $\pi(\mathbf{a}|\mathbf{s})$ from $\arg\max_{\mathbf{a}} Q_\phi(\mathbf{s}, \mathbf{a})$

Important notes:

We can **reuse data** from previous policies!
an **off-policy algorithm** using **replay buffers**

This is **not** a gradient descent algorithm!

Can be readily extended to **multi-task/goal-conditioned** RL

Example: Q-learning Applied to Robotics


1. collect dataset $\{(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)\}$ using some policy

2. set $\mathbf{y}_i \leftarrow r(\mathbf{s}_i, \mathbf{a}_i) + \gamma \max_{\mathbf{a}'_i} Q_\phi(\mathbf{s}'_i, \mathbf{a}'_i)$

3. set $\phi \leftarrow \arg \min_\phi \frac{1}{2} \sum_i \|Q_\phi(\mathbf{s}_i, \mathbf{a}_i) - \mathbf{y}_i\|^2$

Continuous action space?

Simple optimization algorithm ->
Cross Entropy Method (CEM)




1. Start with the normal distribution $N(\mu, \sigma^2)$.

2. Evaluate some parameters from this distribution and select the best (in grey)

3. Compute the mean and std.dev. of the best, add some noise and goto to 1

QT-Opt: Q-learning at Scale




$$\text{minimize } \sum_i (Q(\mathbf{s}_i, \mathbf{a}_i) - [r(\mathbf{s}_i, \mathbf{a}_i) + \max_{\mathbf{a}'_i} Q(\mathbf{s}'_i, \mathbf{a}'_i)])^2$$

Slide adapted from D. Kalashnikov

QT-Opt: Kalashnikov et al. '18, Google Brain

QT-Opt: MDP Definition for Grasping



State: over the shoulder RGB camera image, no depth

Action: 4DOF pose change in Cartesian space + gripper control

Reward: binary reward at the end, if the object was lifted. Sparse. No shaping

Automatic success detection:



QT-Opt: Setup and Results





7 robots collected 580k grasps



Unseen test objects

96% test success rate!



Q-learning

$$\text{Bellman equation: } Q^*(\mathbf{s}_t, \mathbf{a}_t) = \mathbb{E}_{\mathbf{s}' \sim p(\cdot | \mathbf{s}, \mathbf{a})} \left[r(\mathbf{s}, \mathbf{a}) + \gamma \max_{\mathbf{a}'} Q^*(\mathbf{s}', \mathbf{a}') \right]$$

Pros:

- + More sample efficient than on-policy methods
- + Can incorporate off-policy data (including a fully offline setting)
- + Can update the policy even without seeing the reward
- + Relatively easy to parallelize

Cons:

- Harder to apply standard meta-learning algorithms (DP algorithm)
- Lots of “tricks” to make it work
- Potentially could be harder to learn than just a policy

The Plan

Multi-task reinforcement learning problem

Policy gradients & their multi-task/meta counterparts

Q-learning

Multi-task Q-learning

Multi-Task RL Algorithms

Policy: $\pi_\theta(\mathbf{a}|\bar{\mathbf{s}}) \rightarrow \pi_\theta(\mathbf{a}|\bar{\mathbf{s}}, \mathbf{z}_i)$

Q-function: $Q_\phi(\bar{\mathbf{s}}, \mathbf{a}) \rightarrow Q_\phi(\bar{\mathbf{s}}, \mathbf{a}, \mathbf{z}_i)$

Analogous to multi-task supervised learning: stratified sampling, soft/hard weight sharing, etc.

What is different about reinforcement learning?

The data distribution is controlled by the agent!

Why mention it now?

Should we share **data** in addition to sharing **weights**?

An example

Task 1: passing



Task 2: shooting goals




What if you accidentally perform a good pass when trying to shoot a goal?

Store experience as normal. *and* **Relabel** experience with task 2 ID & reward and store.

"**hindsight relabeling**" "hindsight experience replay" (HER)

Goal-conditioned RL with hindsight relabeling

- 1. Collect data $\mathcal{D}_k = \{(\mathbf{s}_{1:T}, \mathbf{a}_{1:T}, \mathbf{s}_g, r_{1:T})\}$ using some policy
- 2. Store data in replay buffer $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_k$
- 3. Perform **hindsight relabeling**:
 - a. Relabel experience in \mathcal{D}_k using last state as goal:
 $\mathcal{D}'_k = \{(\mathbf{s}_{1:T}, \mathbf{a}_{1:T}, \mathbf{s}_T, r'_{1:T})\}$ where $r'_t = -d(\mathbf{s}_t, \mathbf{s}_T)$
 - b. Store relabeled data in replay buffer $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}'_k$
- 4. Update policy using replay buffer \mathcal{D}



< strategies?
the trajectory

Result: exploration challenges alleviated

Why mention it now?

Task 1: close a drawer




Task 2: open a drawer



Can we use episodes from drawer opening task for drawer closing task?

How does that answer change for Q-learning vs Policy Gradient?

Multi-task RL with relabeling

- 
1. Collect data $\mathcal{D}_k = \{(\mathbf{s}_{1:T}, \mathbf{a}_{1:T}, \mathbf{z}_i, r_{1:T})\}$ using some policy
 2. Store data in replay buffer $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}_k$
 3. Perform **hindsight relabeling**:
 - a. Relabel experience in \mathcal{D}_k for task \mathcal{T}_j :
$$\mathcal{D}'_k = \{(\mathbf{s}_{1:T}, \mathbf{a}_{1:T}, \mathbf{z}_j, r'_{1:T})\} \text{ where } r'_t = r_j(\mathbf{s}_t)$$
 - b. Store relabeled data in replay buffer $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}'_k$
 4. Update policy using replay buffer \mathcal{D}
- \leftarrow Which task \mathcal{T}_j to choose?
 - randomly
 - task(s) in which the trajectory gets high reward
- Eysenbach et al. Rewriting History with Inverse RL
Li et al. Generalized Hindsight for RL

When can we apply relabeling?

- reward function form is known, evaluable
- dynamics consistent across goals/tasks
- using an off-policy algorithm*

Hindsight relabeling for goal-conditioned RL

Example: goal-conditioned RL, simulated robot manipulation





Figure 2: Different tasks: *pushing* (top row), *sliding* (middle row) and *pick-and-place* (bottom row). The red ball denotes the goal position.



Time Permitting: What about image observations?

Recall: need a distance function between current and goal state!

$$r'_t = -d(\mathbf{s}_t, \mathbf{s}_T)$$

Use binary 0/1 reward? Sparse, but accurate.

Random, unlabeled interaction is *optimal* under the 0/1 reward of reaching the last state.





Can we use this insight for better learning?

If the data is **optimal**, can we use **supervised imitation learning**?

1. Collect data $\mathcal{D}_k = \{(\mathbf{s}_{1:T}, \mathbf{a}_{1:T})\}$ using some policy
2. Perform **hindsight relabeling**:
 - a. Relabel experience in \mathcal{D}_k using last state as goal:
$$\mathcal{D}'_k = \{(\mathbf{s}_{1:T}, \mathbf{a}_{1:T}, \mathbf{s}_T, r'_{1:T})\} \text{ where } r'_t = -d(\mathbf{s}_t, \mathbf{s}_T)$$
 - b. Store relabeled data in replay buffer $\mathcal{D} \leftarrow \mathcal{D} \cup \mathcal{D}'_k$
3. Update policy using **supervised imitation** on replay buffer \mathcal{D}

Collect data from "human play", perform goal-conditioned imitation.



The Plan

Multi-task reinforcement learning problem

Policy gradients & their multi-task/meta counterparts

Q-learning

Multi-task Q-learning

How **data** can be **shared** across tasks.

Many Remaining Questions: The Next Three Weeks

How can we use a **model** in
multi-task RL?

Model-based RL - Oct 19

What about **meta-RL**
algorithms?

Meta-RL - Oct 21

Can we learn **exploration**
strategies across tasks?

Meta-RL: Learning to explore
- Oct 26

What about **hierarchies** of tasks?

Hierarchical RL - Nov 2

Additional RL Resources

Stanford CS234: Reinforcement Learning

UCL Course from David Silver: Reinforcement Learning

Berkeley CS285: Deep Reinforcement Learning