

# Lifelong Learning

CS 330

# Plan for Today

The lifelong learning **problem statement**

**Basic approaches** to lifelong learning


Can we do **better** than the basics?

Revisiting the problem statement  
from **the meta-learning perspective**

## A brief review of problem statements.

### Multi-Task Learning


Learn to solve a set of tasks.

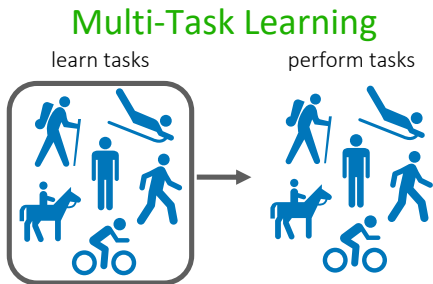


### Meta-Learning


Given i.i.d. task distribution,  
learn a new task efficiently

learn to learn tasks





**Meta-Learning**



In contrast, many real world settings look like:



Our agents may not be given a large batch of data/tasks right off the bat!

Some examples:

- a **student** learning concepts in school
- a deployed **image classification system** learning from a stream of images from users
- a **robot** acquiring an increasingly large set of skills in different environments
- a **virtual assistant** learning to help different users with different tasks at different points in time
- a **doctor's assistant** aiding in medical decision-making

# Some Terminology

## Sequential learning settings

online learning, lifelong learning, continual learning, incremental learning, streaming data

distinct from **sequence data** and **sequential decision-making**

# What is the lifelong learning *problem statement*?

## Exercise:

1. Pick an example setting.
2. Discuss problem statement in your break-out room:
  - (a) how would you set-up an experiment to develop & test your algorithm?
  - (b) what are desirable/required properties of the algorithm?
  - (c) how do you evaluate such a system?

## Example settings:

- 
- A. a **student** learning concepts in school
  - B. a deployed **image classification system** learning from a stream of images from users
  - C. a **robot** acquiring an increasingly large set of skills in different environments
  - D. a **virtual assistant** learning to help different users with different tasks at different points in time
  - E. a **doctor's assistant** aiding in medical decision-making

Desirable properties/considerations	Evaluation setup

# What is the lifelong learning *problem statement*?

Problem variations:

- **task/data order**: i.i.d. vs. predictable vs. curriculum vs. adversarial
- **discrete** task boundaries vs. **continuous** shifts (vs. both)
- **known** task boundaries/shifts vs. **unknown**

Some considerations:

- model **performance**
- data **efficiency**
- **computational** resources
- **memory**
- others: privacy, interpretability, fairness,  
test time compute & memory

Substantial variety in problem statement!



# What is the lifelong learning *problem statement*?

General [supervised] online learning problem:

for  $t = 1, \dots, n$

observe  $x_t$

<— if **observable task boundaries**: observe  $x_t, z_t$

predict  $\hat{y}_t$

observe label  $y_t$

**i.i.d. setting:**  $x_t \sim p(x), y_t \sim p(y|x)$

$p$  not a function of  $t$

otherwise:  $x_t \sim p_t(x), y_t \sim p_t(y|x)$

**streaming setting:** cannot store  $(x_t, y_t)$

- lack of memory
- lack of computational resources
- privacy considerations
- want to study neural memory mechanisms

true in some cases, but not in many cases!

9 - recall: replay buffers

# What do you want from your lifelong learning algorithm?

**minimal regret** (that grows slowly with  $t$ )

**regret**: cumulative loss of learner — cumulative loss of best learner in hindsight

$$\text{Regret}_T := \sum_1^T \mathcal{L}_t(\theta_t) - \min_{\theta} \sum_1^T \mathcal{L}_t(\theta)$$

(cannot be evaluated in practice, useful for analysis)

Regret that grows linearly in  $t$  is trivial. Why?

# What do you want from your lifelong learning algorithm?

## **positive & negative transfer**

positive **forward** transfer: previous tasks cause you to do better on future tasks  
compared to learning future tasks from scratch

positive **backward** transfer: current tasks cause you to do better on previous tasks  
compared to learning past tasks from scratch

positive -> negative : better -> worse

# Plan for Today

The lifelong learning **problem statement**

**Basic approaches** to lifelong learning

Can we do **better** than the basics?

Revisiting the problem statement  
from **the meta-learning perspective**

# Approaches

Store all the data you've seen so far, and train on it. —> **follow the leader** algorithm

- + will achieve very strong performance

- **computation intensive** —> Continuous fine-tuning can help.

- **can be memory intensive** [depends on the application]

Take a gradient step on the datapoint you observe. —> **stochastic gradient descent**

- + computationally cheap

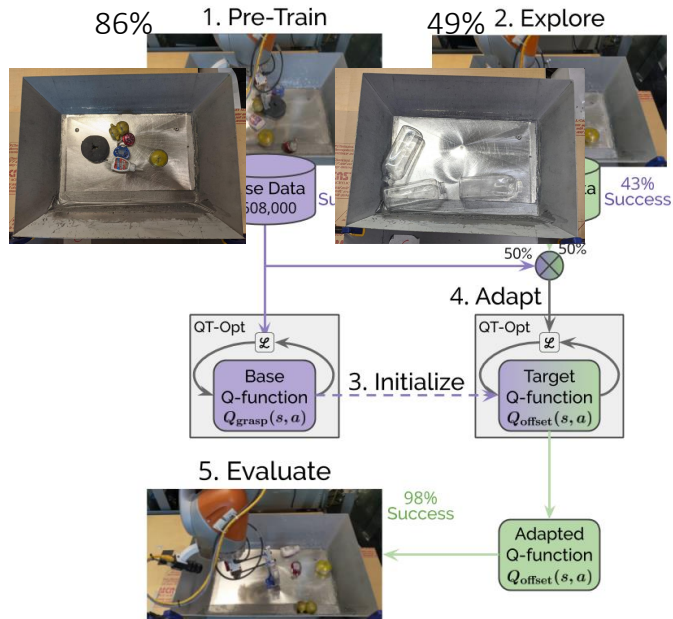
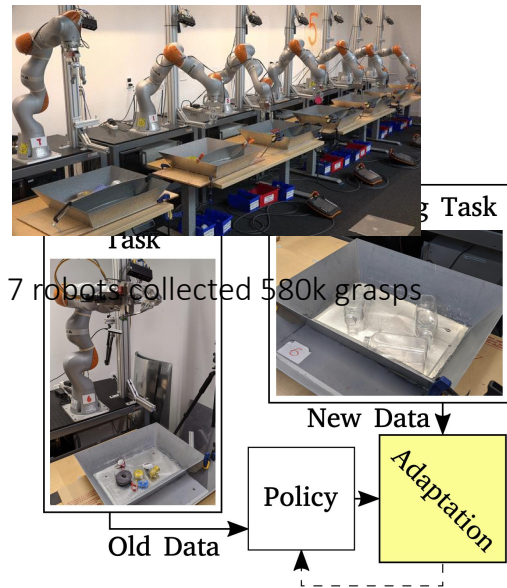
- + requires 0 memory

- subject to **negative backward transfer**  
“forgetting”

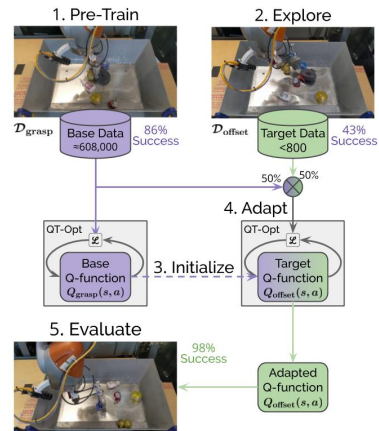
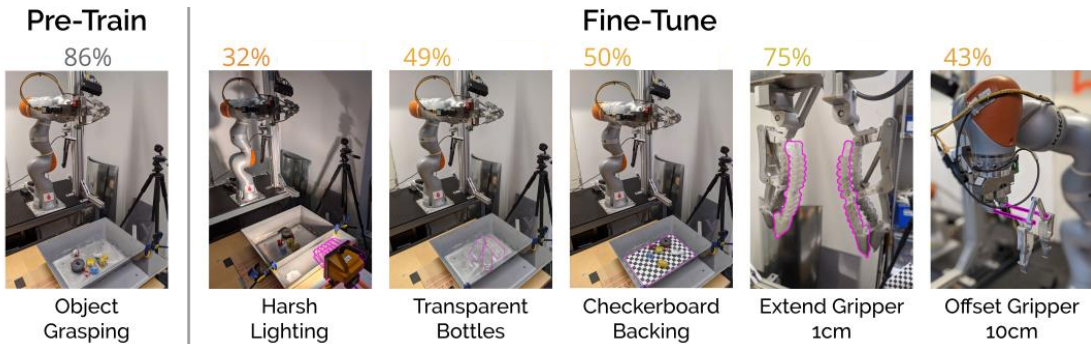
sometimes referred to as  
catastrophic forgetting

- slow learning

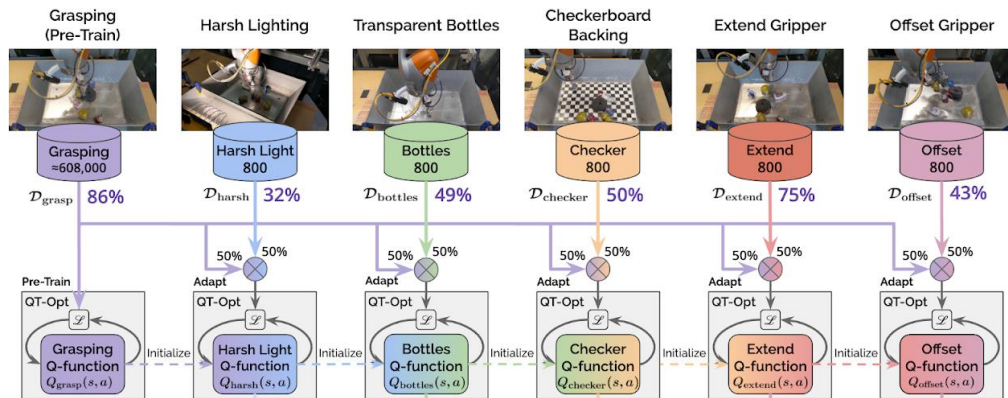
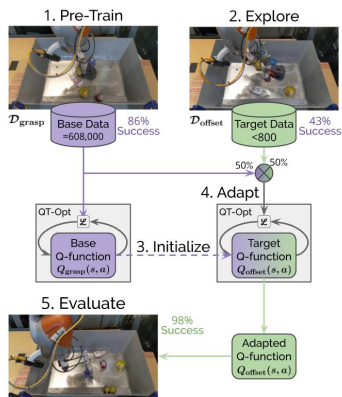
# Very simple continual RL algorithm



# Very simple continual RL algorithm



# Very simple continual RL algorithm



Can we do better?



# Plan for Today

The lifelong learning **problem statement**

**Basic approaches** to lifelong learning

Can we do **better** than the basics?

Revisiting the problem statement  
from **the meta-learning perspective**

Case Study: Can we modify vanilla SGD to avoid negative backward transfer?  
(from scratch)

- Idea:
- (1) store small amount of data per task in **memory**
  - (2) when making updates for new tasks, ensure that they don't **unlearn** previous tasks

How do we accomplish (2)?

learning predictor  $y_t = f_\theta(x_t, z_t)$     memory:  $\mathcal{M}_k$  for task  $z_k$

For  $t = 0, \dots, T$

minimize  $\mathcal{L}(f_\theta(\cdot, z_t), (x_t, y_t))$

subject to  $\mathcal{L}(f_\theta, \mathcal{M}_k) \leq \mathcal{L}(f_\theta^{t-1}, \mathcal{M}_k)$  for all  $z_k < z_t$

(i.e. s.t. loss on previous tasks doesn't get worse)

Assume local  
linearity:

$$\langle g_t, g_k \rangle := \left\langle \frac{\partial \mathcal{L}(f_\theta, (x_t, y_t))}{\partial \theta}, \frac{\mathcal{L}(f_\theta, \mathcal{M}_k)}{\partial \theta} \right\rangle \geq 0 \quad \text{for all } z_k < z_t$$

Can formulate & solve as a QP.







# Experiments

Problems:

- MNIST permutations
- MNIST rotations
- CIFAR-100 (5 new classes/task)

BWT: backward transfer,  
FWT: forward transfer

Total memory size:  
5012 examples



If we take a step back... do these experimental domains make sense?

Can we meta-learn how to avoid negative backward transfer?

Javed & White. *Meta-Learning Representations for Continual Learning*. NeurIPS '19

Beaulieu et al. *Learning to Continually Learn*. '20

# Plan for Today

The lifelong learning **problem statement**

**Basic approaches** to lifelong learning

Can we do **better** than the basics?


Revisiting the problem statement  
from **the meta-learning perspective**

# What might be wrong with the online learning formulation?


## Online Learning

(Hannan '57, Zinkevich '03)

Perform sequence of tasks  
while minimizing static regret.



More realistically:




# What might be wrong with the online learning formulation?

## Online Learning


(Hannan '57, Zinkevich '03)

Perform sequence of tasks while minimizing static regret.



## Online Meta-Learning

Efficiently learn a sequence of tasks from a non-stationary distribution.



Primarily a difference in *evaluation*, rather than the *data stream*.



## The Online Meta-Learning Setting

for task  $t = 1, \dots, n$

observe  $\mathcal{D}_t^{tr}$

use update procedure  $\Phi(\theta_t, \mathcal{D}_t^{tr})$  to produce parameters  $\phi_t$

observe  $x_t$

predict  $\hat{y}_t = f_{\phi_t}(x_t)$

observe label  $y_t$

Standard online learning setting

**Goal:** Learning algorithm with sub-linear

$$\text{Regret}_T := \sum_{t=1}^T \ell_t(\Phi_t(\theta_t)) - \min_{\theta \in \Theta} \sum_{t=1}^T \ell_t(\Phi_t(\theta))$$

Loss of algorithm

Loss of best algorithm in hindsight

# Can we apply meta-learning in lifelong learning settings?

Recall the **follow the leader** (FTL) algorithm:

Store all the data you've seen so far, and train on it.

Deploy model on current task.

Follow the *meta*-leader (FTML) algorithm:

Store all the data you've seen so far, and **meta**-train on it.


Run update procedure on the current task.

What meta-learning algorithms are well-suited for FTML?

What if  $p_t(\mathcal{T})$  is non-stationary?

## Optimization vs. Black-Box Adaptation

How well can learning procedures generalize to similar, but extrapolated tasks?



Does this structure come at a cost?


Finn & Levine ICLR 18

# Experiments

Experiment with **sequences of tasks**:

- Colored, rotated, scaled **MNIST**
- **3D object pose prediction**
- **CIFAR-100** classification

Example pose prediction tasks



plane


car

chair

# Experiments


- Comparisons:
- **TOE** (train on everything): train on all data so far
  - **FTL** (follow the leader): train on all data so far, fine-tune on current task
  - **From Scratch**: train from scratch on each task

Learning efficiency



Task index

Learning proficiency (error)



Task index

**Follow The Meta-Leader**  
learns each new task faster & with greater proficiency,  
approaches **few-shot learning** regime

# Takeaways

Many flavors of lifelong learning, all under the same name.

Defining the problem statement is often the hardest part

Meta-learning can be viewed as a slice of the lifelong learning problem.

A very open area of research.