

Tensorflow 2.x Review Session

CS330: Deep Multi-task and Meta Learning

9/17/2019

Rafael Rafailov

Overview

1. Installation

- a. Installing on your machine
- b. Using GPUs
- c. Using Google Colab

2. Tensorflow Basics

- a. Data pipelines
- b. Autograd in TF 2.0
- c. Models
- d. Optimizers
- e. Training loop

3. Other topics

- a. Layers with memory (for HW1)
- b. Tensorflow Probability

Installing on your machine/cloud instance

Directly:

```
# Requires the latest pip
```

```
pip install --upgrade pip
```

```
# Current stable release for CPU and GPU
```

```
pip install tensorflow(-gpu)
```

Important - make sure you're actually using the GPU

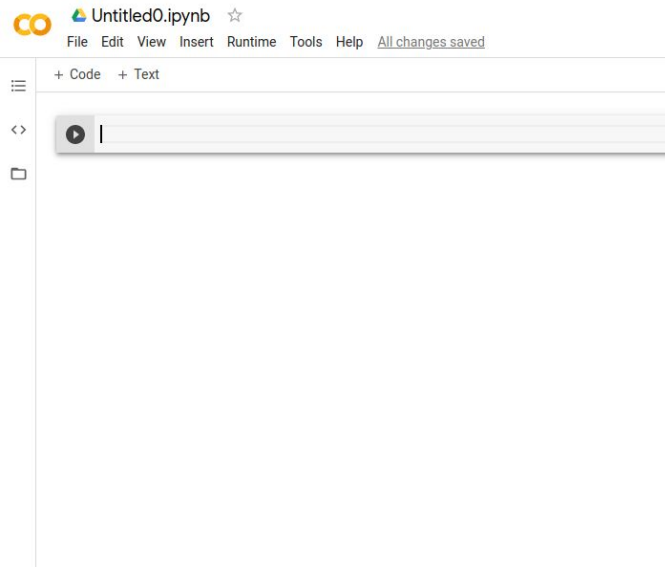
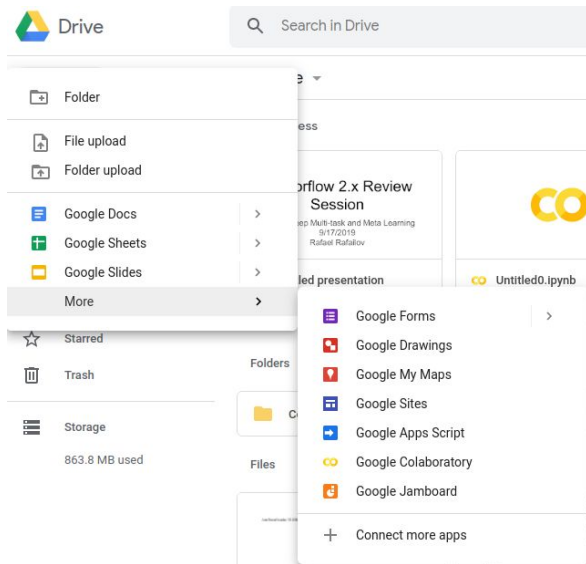
```
tf.test.is_gpu_available(  
    cuda_only=False, min_cuda_compute_capability=None  
)
```

```
>>> tf.test.is_gpu_available()  
WARNING:tensorflow:From <stdin>:1: is_gpu_available (from tensorflow.python.framework.test_util) is deprecated and will be removed in a future version.  
Instructions for updating:  
Use 'tf.config.list_physical_devices('GPU')' instead.  
2020-09-14 09:24:25.470684: I tensorflow/core/platform/cpu_feature_guard.cc:143] Your CPU supports instructions that this TensorFlow binary was not compiled to use: AVX2 FMA  
2020-09-14 09:24:25.516884: I tensorflow/core/platform/profile_utils/cpu_utils.cc:102] CPU Frequency: 2599990000 Hz  
2020-09-14 09:24:25.519302: I tensorflow/compiler/xla/service/service.cc:168] XLA service 0x7f26d800b20 initialized for platform Host (this does not guarantee that XLA will be used). Devices:  
2020-09-14 09:24:25.519417: I tensorflow/compiler/xla/service/service.cc:176] StreamExecutor device (0): Host, Default Version  
2020-09-14 09:24:25.528671: I tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully opened dynamic library libcudart.so.1  
2020-09-14 09:24:25.634104: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:981] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so r  
returning NUMA node zero  
2020-09-14 09:24:25.634540: I tensorflow/compiler/xla/service/service.cc:168] XLA service 0x55ebd2f094c0 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:  
2020-09-14 09:24:25.634560: I tensorflow/compiler/xla/service/service.cc:176] StreamExecutor device (0): GeForce RTX 2080 with Max-Q Design, Compute Capability 7.5  
2020-09-14 09:24:25.635221: I tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:981] successful NUMA node read from SysFS had negative value (-1), but there must be at least one NUMA node, so r  
returning NUMA node zero  
2020-09-14 09:24:25.635517: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1561] Found device 0 with properties:  
pciBusID: 0000:01:00.0 name: GeForce RTX 2080 with Max-Q Design computeCapability: 7.5  
coreClock: 1.095GHz coreCount: 46 deviceMemorySize: 7.79GiB deviceMemoryBandwidth: 357.69GiB/s  
2020-09-14 09:24:25.636738: I tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully opened dynamic library libcudart.so.10.1  
2020-09-14 09:24:25.636852: W tensorflow/stream_executor/platform/default/dso_loader.cc:55] Could not load dynamic library 'libcublas.so.10'; derror: libcublas.so.10: cannot open shared object fi  
le: No such file or directory  
2020-09-14 09:24:25.636905: W tensorflow/stream_executor/platform/default/dso_loader.cc:55] Could not load dynamic library 'libcufft.so.10'; derror: libcufft.so.10: cannot open shared object file  
: No such file or directory  
2020-09-14 09:24:25.636955: W tensorflow/stream_executor/platform/default/dso_loader.cc:55] Could not load dynamic library 'libcurand.so.10'; derror: libcurand.so.10: cannot open shared object fi  
le: No such file or directory  
2020-09-14 09:24:25.637804: W tensorflow/stream_executor/platform/default/dso_loader.cc:55] Could not load dynamic library 'libcusolver.so.10'; derror: libcusolver.so.10: cannot open shared objec  
t file: No such file or directory  
2020-09-14 09:24:25.637853: W tensorflow/stream_executor/platform/default/dso_loader.cc:55] Could not load dynamic library 'libcusparse.so.10'; derror: libcusparse.so.10: cannot open shared objec  
t file: No such file or directory  
2020-09-14 09:24:25.661940: I tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully opened dynamic library libcudnn.so.7  
2020-09-14 09:24:25.661967: W tensorflow/core/common_runtime/gpu/gpu_device.cc:1598] Cannot dlopen some GPU libraries. Please make sure the missing libraries mentioned above are installed properly  
if you would like to use GPU. Follow the guide at https://www.tensorflow.org/install/gpu for how to download and setup the required libraries for your platform.  
Skipping registering GPU devices...  
2020-09-14 09:24:25.661985: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1102] Device interconnect StreamExecutor with strength 1 edge matrix:  
2020-09-14 09:24:25.661991: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1108] 0  
2020-09-14 09:24:25.661998: I tensorflow/core/common_runtime/gpu/gpu_device.cc:1121] 0: N  
False
```

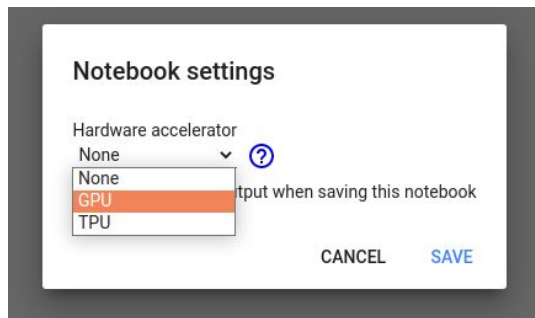
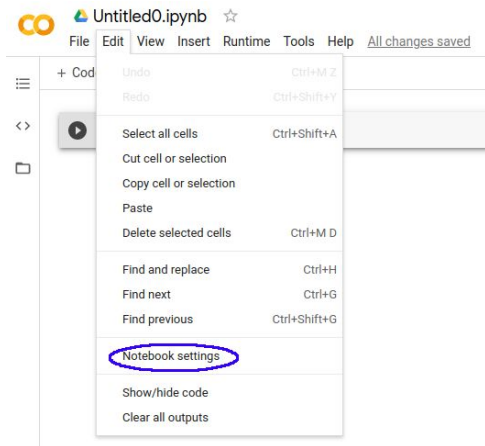
Need to have CUDA packages installed

1. Instructions are here - <https://www.tensorflow.org/install/gpu>
2. Alternatively, to install CUDA dependencies using conda (I find this easier, especially if you do not have sudo access on the machine):
<https://towardsdatascience.com/managing-cuda-dependencies-with-conda-89c5d817e7e1>
3. To check CUDA version - `nvidia-smi`

Using Google Colaboratory Notebooks



Using GPU in Colabs - no GPU by default!



```
[1] import tensorflow as tf
```

```
tf.test.is_gpu_available()
```

```
WARNING:tensorflow:From <ipython-input-2-17bb7203622b>:1:  
Instructions for updating:  
Use `tf.config.list_physical_devices('GPU')` instead.  
True
```

Getting started

1. Colab notebooks already have tensorflow (and GPUs) set up.
2. Homeworks should be doable on CPUs too, but might take a bit longer (couple of hours).
3. Only need to set-up TF 2.x if you're using a separate instance for your project (can be done in PyTorch too).

Tensorflow Data pipelines

```
dataset = tf.data.Dataset.from_generator(generator, types, shapes)
```

```
dataset = dataset.batch(batch_size, drop_remainder=True)
```

```
dataset = dataset.map(preprocess)
```

```
dataset = dataset.prefetch(10)
```

1. Create a TF dataset object from python generator object
2. Create a batched dataset (i.e. sample batches of batch_size)
3. Apply preprocess() to each batch before returning it (e.g. normalize images to [0,1])
4. Preload batches while computation is running - significant speed up in data I/O

More functionality: https://www.tensorflow.org/api_docs/python/tf/data/Dataset

Tensorflow Gradients and Autodiff

```
w = tf.Variable(tf.random.normal((3, 2)), name='w')
b = tf.Variable(tf.zeros(2, dtype=tf.float32), name='b')
x = [[1., 2., 3.]]
```

```
with tf.GradientTape(persistent=True) as tape:
    y = x @ w + b
    loss = tf.reduce_mean(y**2)
```

-> Gradient tape tracks differentiable operations
-> persistent = True keeps compute graph after tape.gradient

```
[dl_dw, dl_db] = tape.gradient(loss, [w, b])
```

-> Computes tracked variable grads

```
print(y)
print(dl_db)
tf.Tensor([[ 1.9099498 -8.337775 ]], shape=(1, 2), dtype=float32)
tf.Tensor([ 1.9099499 -8.337775 ], shape=(2,), dtype=float32)
```

Tensorflow Gradients and Autodiff

```
x0 = tf.Variable(3.0, name='x0')  
x1 = tf.Variable(3.0, name='x1', trainable=False)  
x2 = tf.Variable(2.0, name='x2') + 1.0  
x3 = tf.constant(3.0, name='x3')
```

```
with tf.GradientTape() as tape:  
    y = (x0**2) + (x1**2) + (x2**2)
```

```
grad = tape.gradient(y, [x0, x1, x2, x3])
```

```
tf.Tensor(6.0, shape=(), dtype=float32)
```

```
None
```

```
None
```

```
None
```

Tensorflow Models

```
import tensorflow as tf
from tensorflow.keras.layers import Dense, Flatten, Conv2D
from tensorflow.keras import Model

class MyModel(Model):                                -> Define model
    def __init__(self):
        super(MyModel, self).__init__()
        self.conv1 = Conv2D(32, 3, activation='relu')
        self.flatten = Flatten()
        self.d1 = Dense(128, activation='relu')        -> Define model layers
        self.d2 = Dense(10)

    def call(self, x):
        x = self.conv1(x)
        x = self.flatten(x)                            -> Model processing
        x = self.d1(x)
        return self.d2(x)

# Create an instance of the model
model = MyModel()                                    -> Initialize Model
```

TF Keras layers: https://www.tensorflow.org/api_docs/python/tf/keras/layers

Tensorflow Losses and Metrics

Losses :

```
loss_object = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)  
loss_object = tf.keras.losses.MeanSquaredError()
```

Other losses: https://www.tensorflow.org/api_docs/python/tf/keras/losses

Metrics:

```
test_loss = tf.keras.metrics.Mean()  
test_accuracy = tf.keras.metrics.SparseCategoricalAccuracy()  
test_top_k_accuracy = tf.keras.metrics.SparseTopK_categoricalAccuracy(k=5)
```

Other metrics: https://www.tensorflow.org/api_docs/python/tf/keras/metrics

Tensorflow Optimizers

```
optimizer = tf.keras.optimizers.Adam(  
    learning_rate=0.001, beta_1=0.9, beta_2=0.999, epsilon=1e-07, amsgrad=False,  
    name='Adam')
```

```
optimizer = tf.keras.optimizers.RMSprop(  
    learning_rate=0.001, rho=0.9, momentum=0.0, epsilon=1e-07, centered=False,  
    name='RMSprop')
```

```
optimizer = tf.keras.optimizers.SGD(  
    learning_rate=0.01, momentum=0.0, nesterov=False, name='SGD')
```

Available optimizers: https://www.tensorflow.org/api_docs/python/tf/keras/optimizers

Putting it all together

```
def train_step(images, labels):  
    with tf.GradientTape() as tape:  
        predictions = model(images, training=True)  
        loss = loss_object(labels, predictions)  
    gradients = tape.gradient(loss, model.trainable_variables)  
    optimizer.apply_gradients(zip(gradients, model.trainable_variables))
```

-> sets behaviour for
Dropout() layers etc
-> compute model grads
-> apply grads to weights

Applying gradients by hand

Applying gradients by hand (useful for optimization-based meta-learning):

```
gradients = tape.gradient(inner_loss, trainable_weights)
new_weights = ([weight - lr_inner * grad for weight, grad in zip(trainable_weights,
                                                                  gradients)])
```

Important!

```
model.set_weights(weights)
variable.assign(weight)
```

Will break the computation graph (will become clear later on)!

Let's run it!

Colab is here:

<https://colab.research.google.com/github/tensorflow/docs/blob/master/site/en/tutorials/quickstart/advanced.ipynb>

Recurrent Cells

```
tf.keras.layers.LSTMCell(units)
```

```
import tensorflow as tf
```

```
inputs = tf.random.normal([32, 10, 8])
```

```
cell = tf.keras.layers.LSTMCell(4)
```

```
state = cell.get_initial_state(batch_size = 32, dtype = tf.float32)
```

```
output, state = cell(inputs[:,0], state)
```

```
print(output.shape)
```

```
print(state[0].shape)
```

```
print(state[1].shape)
```

```
(32, 4)
```


```
(32, 4)
```

```
(32, 4)
```

-> batch x length x size of data

-> initialize cell state

-> process data one at a time



Recurrent Networks

```
inputs = tf.random.normal([32, 10, 8])
rnn = tf.keras.layers.RNN(tf.keras.layers.LSTMCell(4))    -> wrap cell to process sequence
output = rnn(inputs)
print(output.shape)
(32, 4)
rnn = tf.keras.layers.RNN(
    tf.keras.layers.LSTMCell(4),
    return_sequences=True,
    return_state=True)
whole_seq_output, final_memory_state, final_carry_state = rnn(inputs)
print(whole_seq_output.shape)
(32, 10, 4)
print(final_memory_state.shape)
(32, 4)
print(final_carry_state.shape)
(32, 4)
```

Recurrent Networks

```
inputs = tf.random.normal([32, 10, 8])
lstm = tf.keras.layers.LSTM(4)
output = lstm(inputs)
print(output.shape)
(32, 4)
lstm = tf.keras.layers.LSTM(4, return_sequences=True, return_state=True)
whole_seq_output, final_memory_state, final_carry_state = lstm(inputs)
print(whole_seq_output.shape)
(32, 10, 4)
print(final_memory_state.shape)
(32, 4)
print(final_carry_state.shape)
(32, 4)
```

Black-box model for HW1

Tensorflow Probability

Installation:

```
pip install --upgrade tensorflow-probability
```

Uses:

1. Generative models (i.e. VAEs, Autoregressive Models, Normalizing Flows)
2. Statistical Models (i.e. Bayesian Models, Hamiltonian MCMC)
3. Reinforcement Learning (i.e. stochastic policies)

Some advanced examples:

https://github.com/tensorflow/probability/tree/master/tensorflow_probability/examples

Tensorflow Probability

```
import tensorflow as tf
import tensorflow_probability as tfp

mean = tf.Variable([1.0, 2.0, 3.], name='mean')
std = tf.Variable([0.1, 0.1, 0.1], name='std')
var = tf.constant([3.0, 0.1, 2.0], name='var')

with tf.GradientTape(persistent=True) as tape:
    dist = tfp.distributions.Normal(loc = mean, scale = std)
    s = dist.sample()
    loss1 = tf.reduce_mean(s**2)
    loss2 = tf.reduce_mean(dist.log_prob(var))
    loss3 = tf.reduce_mean(dist.log_prob(s))

grad1 = tape.gradient(loss1, [mean])
grad2 = tape.gradient(loss2, [mean])
grad3 = tape.gradient(loss3, [mean])

print(grad1)
print(grad2)
print(grad3)

[<tf.Tensor: shape=(3,), dtype=float32, numpy=array([0.63096106, 1.3687671 , 1.9575679 ], dtype=float32)>]
[<tf.Tensor: shape=(3,), dtype=float32, numpy=array([ 66.66667 , -63.333336, -33.333336], dtype=float32)>]
[<tf.Tensor: shape=(3,), dtype=float32, numpy=array([0., 0., 0.], dtype=float32)>]
```

TF Agents

```
pip install tf-agents
```

```
import tensorflow as tf
from tf_agents.networks import q_network
from tf_agents.agents.dqn import dqn_agent
```

```
q_net = q_network.QNetwork(
    train_env.observation_spec(),
    train_env.action_spec(),
    fc_layer_params=(100,))

agent = dqn_agent.DqnAgent(
    train_env.time_step_spec(),
    train_env.action_spec(),
    q_network=q_net,
    optimizer=optimizer,
    td_errors_loss_fn=common.element_wise_squared_loss,
    train_step_counter=tf.Variable(0))
```

```
agent.initialize()
```

<https://www.tensorflow.org/agents>

Questions?